

Γ

# Practical AI

L

## *Part 1- Intelligent Search*

### Search

(Intelligent) Search

Constraint Satisfaction

Adversarial Search (Games)



Universität  
Zürich<sup>UZH</sup>



Dynamic and Distributed  
Information Systems



Note: Slides based on various slide sets for the Russel and Norvig AI book



# Overview – Part 1



- **Problem Definition?**
- Search
  - Breadth first, Depth first, Iterative deepening
- Informed (or Intelligent) Search
  - Best First
  - A\*, Hill climbing, Simulated annealing
- Constraint Satisfaction
  - Forward checking, Constraint propagation
- Adversarial Search (Games)

# Planning

- Input
  - Description of set of all possible **states** of the world (in some knowledge representation language)
  - Description of initial state of world
  - Description of goal
  - Description of available actions
    - May include **costs** for performing actions
- Output
  - **Sequence** of actions that convert the initial state into one that satisfies the goal
  - May wish to minimize length or cost of plan



# Example problem: Traveling in Romania

## ***Situation:***

- On holiday in Romania; currently in Arad
- Flight leaves tomorrow from Bucharest

## ***Formulate Goal:***

- be in Bucharest

## ***Formulate Problem:***

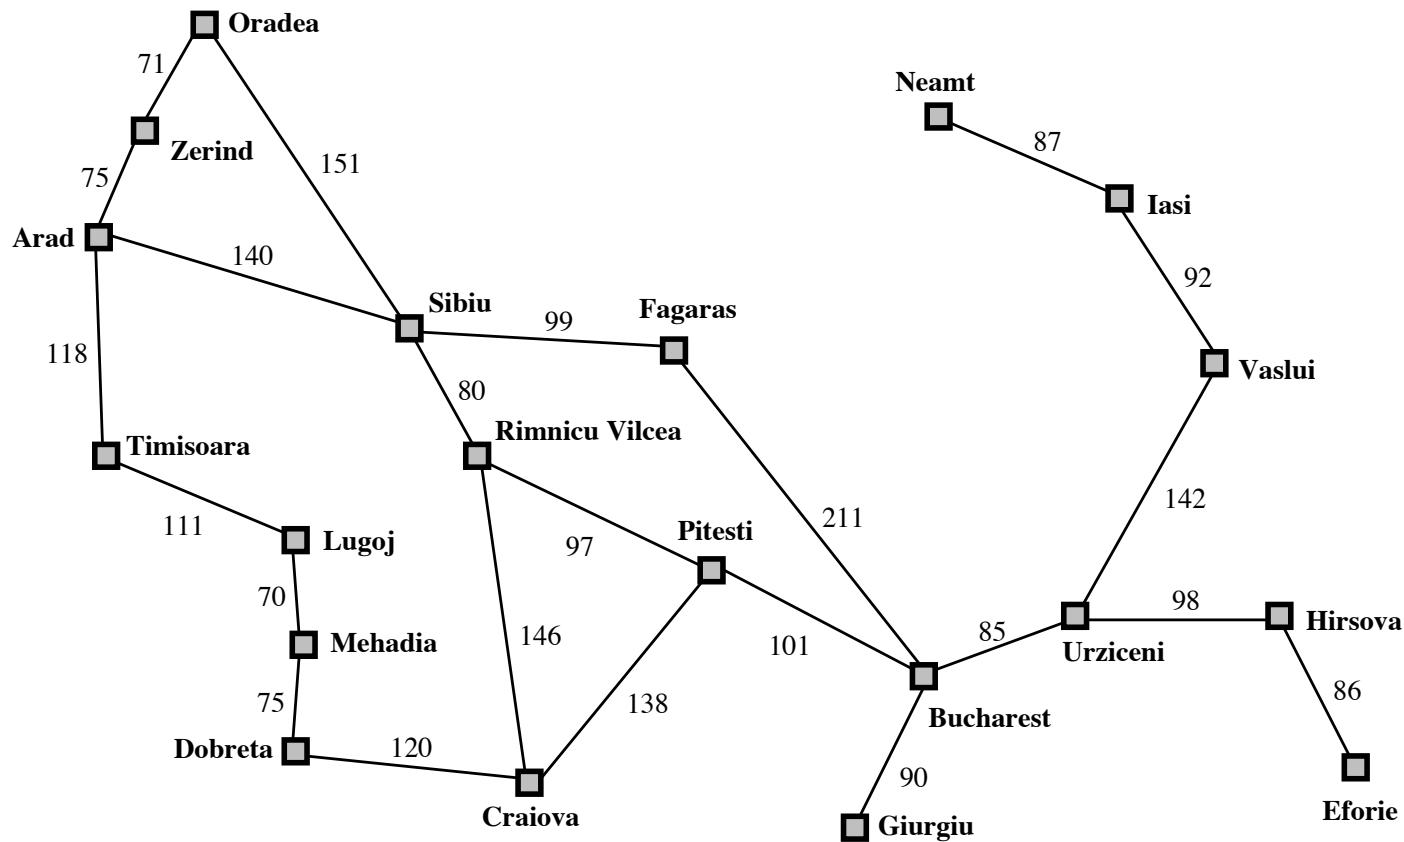
- states: various cities
- actions: drive between cities

## ***Find solution:***

- sequence of cities  
e.g., Arad, Sibiu, Fagaras, Bucharest



# Romania





# Problem Types



Deterministic, fully observable

→ single-state problem

- Agent knows exactly which state it will be in; solution is a sequence

Non-observable

→ conformant problem

- Agent may have no idea where it is; solution (if any) is a sequence

Nondeterministic and/or partially observable → contingency problem

- percepts provide *new* information about current state
- solution is a *tree* or *policy*
- often *interleave* search, execution

Unknown state space

→ exploration problem (online)



# Problem Formulation



A **problem** is defined by four items:

- **initial state**: e.g., ``at Arad''
- **successor function**:  $S(x)$  =set of action-state pairs
  - e.g.,  $S(\text{Arad}) = \{\langle \text{Arad} \rightarrow \text{Zerind}, \text{Zerind} \rangle, \dots\}$
- **goal test**, can be
  - *explicit*, e.g.,  $x = \text{"at Bucharest"}$
  - *implicit*, e.g.,  $\text{NoDirt}(x)$
- **path cost** (additive)
  - e.g., sum of distances, number of actions executed, etc.
  - $c(x,a,y)$  is the *step cost*, assumed to be  $\geq 0$
- A **solution** is a sequence of actions leading from the initial state to a goal state

# Defining the Problem

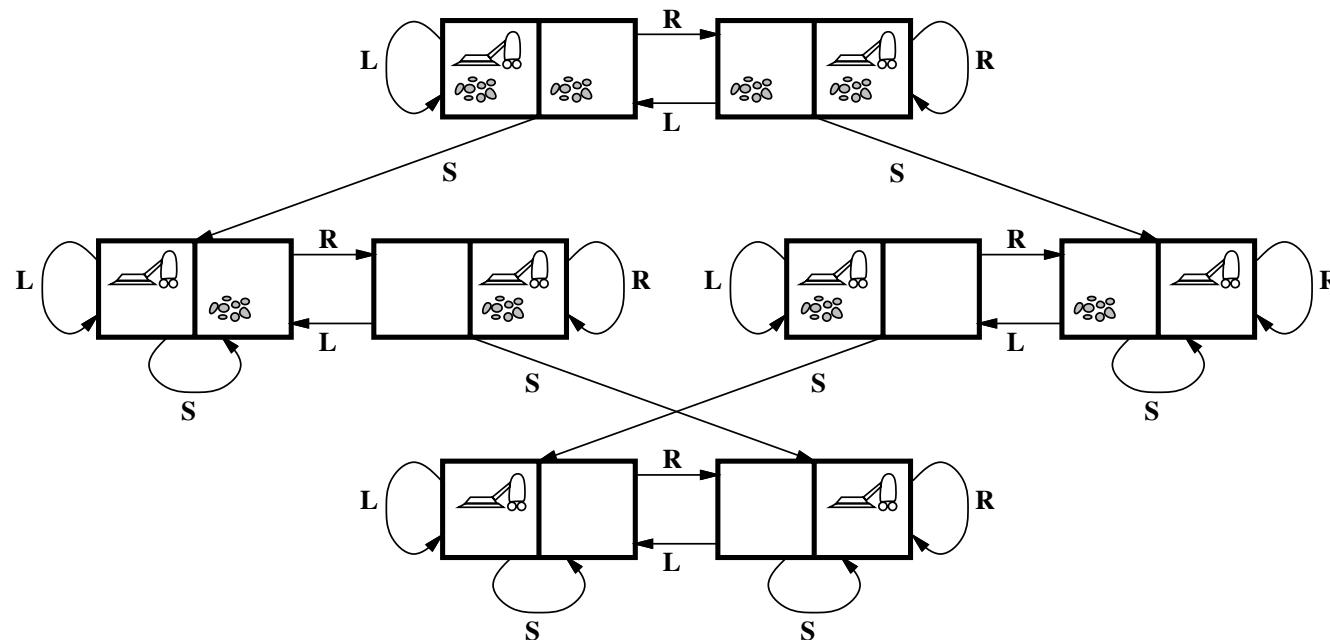
## *Example: Vacuum World*

**States?**

**Actions?**

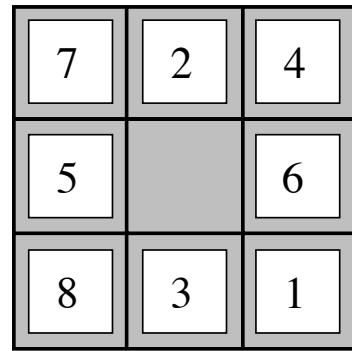
**Goal test?**

**Path cost?**

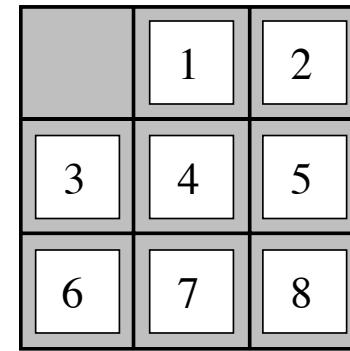


# Defining the Problem

## *Example: The 8-Puzzle*



Start State



Goal State

**States?**

**Actions?**

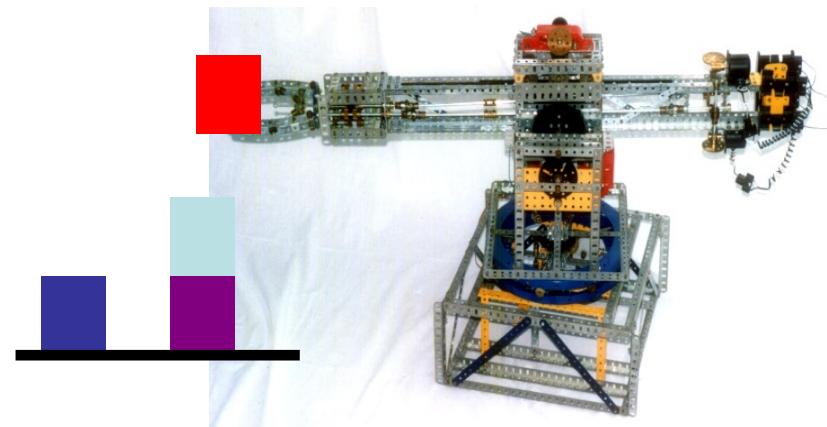
**Goal test?**

**Path cost?**

# Defining the Problem

## *Example: Robot Control - Blocks World*

- Input:
  - **States?**
  - **Actions?**
  - **Goal test?**
  - **Path cost?**
- Output:
  - **Action Plan?**





# STRIPS Representation

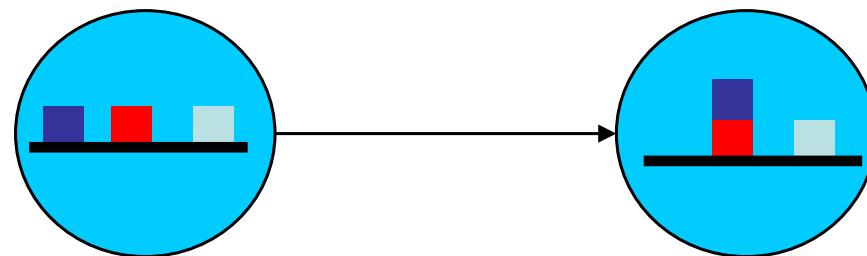


- Description of initial state of world
  - Set of *propositions* that completely describes a world
  - { (block a) (block b) (block c) (on-table a)  
(on-table b) (clear a) (clear b) (clear c)  
(arm-empty) }
- Description of goal (i.e. set of desired worlds)
  - Set of propositions that *partially* describes a world
  - { (on a b) (on b c) }
- Description of available actions

# └ STRIPS Representation

## └ How Represent Actions?

- World = set of propositions true in that world
- Actions:
  - Precondition: conjunction of propositions
  - Effects: propositions made true & propositions made false (deleted from the state description)



*operator: stack\_B\_on\_R*

*precondition: (on B Table) (clear R)*

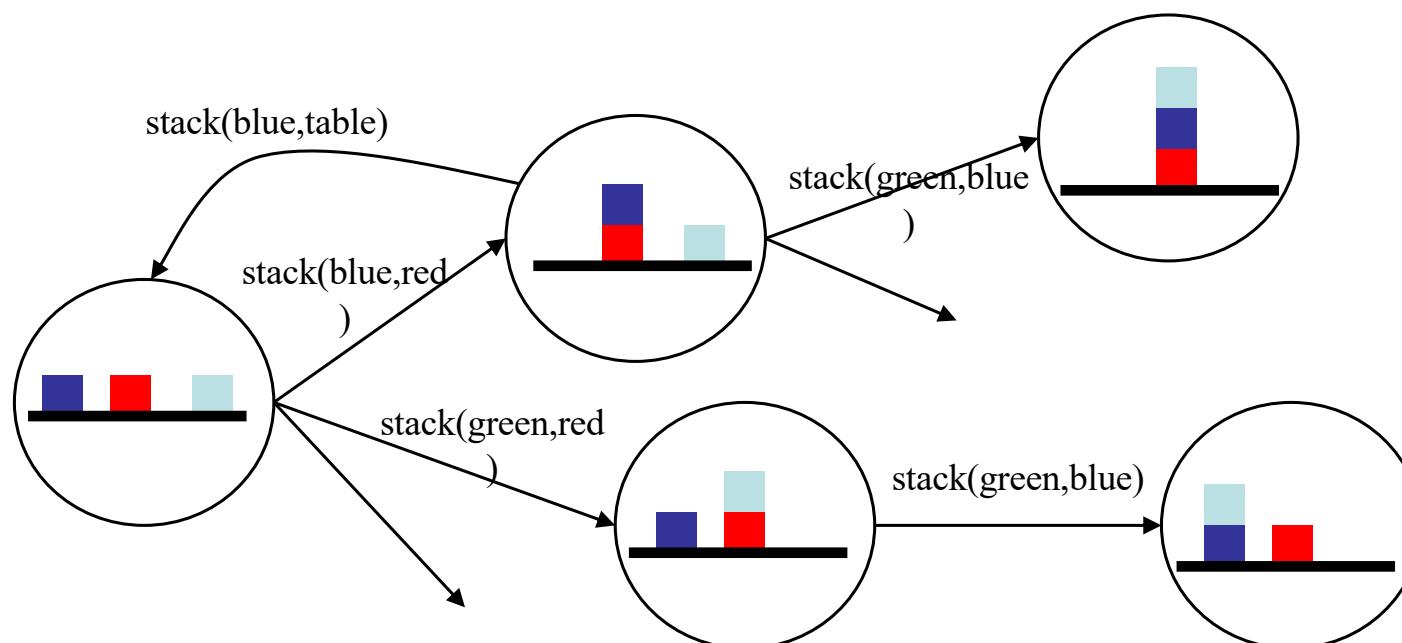
*effect: (on B R) (:not (clear R))*



## Planning as graph search



- Planning can be viewed as finding paths in a graph, where the graph is implicitly specified by the set of actions
- **Blocks world:**
  - vertex = relative positions of all blocks
  - edge = robot arm stacks one block





# Overview – Part 1



- Problem Definition?
- **Search**
  - Breadth first, Depth first, Iterative deepening
- Informed (or Intelligent) Search
  - Best First
  - A\*, Hill climbing, Simulated annealing
- Constraint Satisfaction
  - Forward checking, Constraint propagation
- Adversarial Search (Games)



## Searching: Tree Search

- Basic idea: offline, simulated exploration of state space by generating successors of already-explored states  
(a.k.a. expanding states)

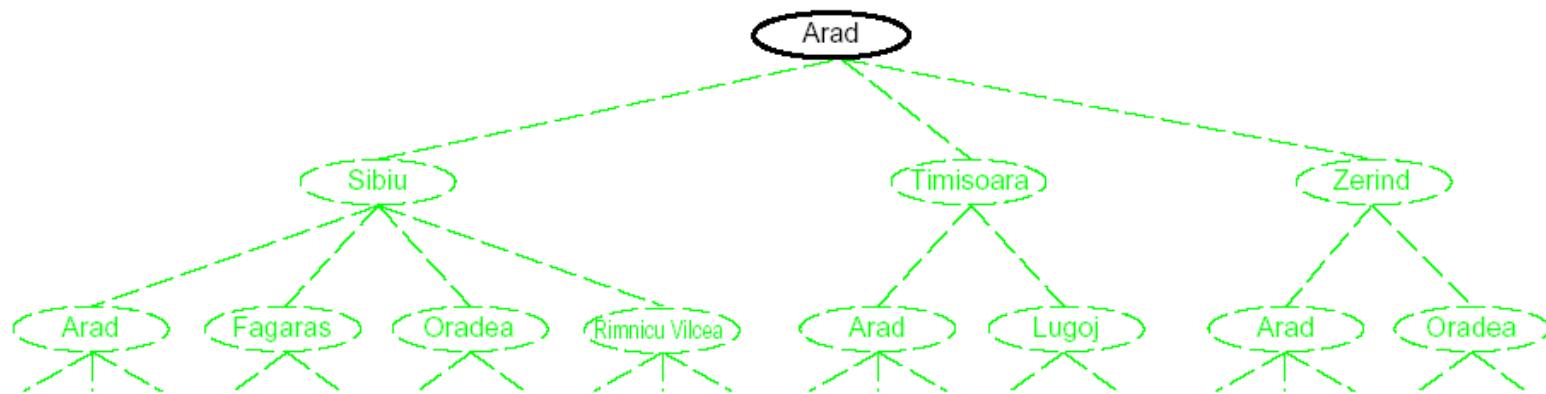
---

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
    initialize the search tree using the initial state of problem
    loop do
        if there are no candidates for expansion then return failure
        choose a leaf node for expansion according to strategy
        if the node contains a goal state then return the corresponding solution
        else expand the node and add the resulting nodes to the search tree
    end
```

---

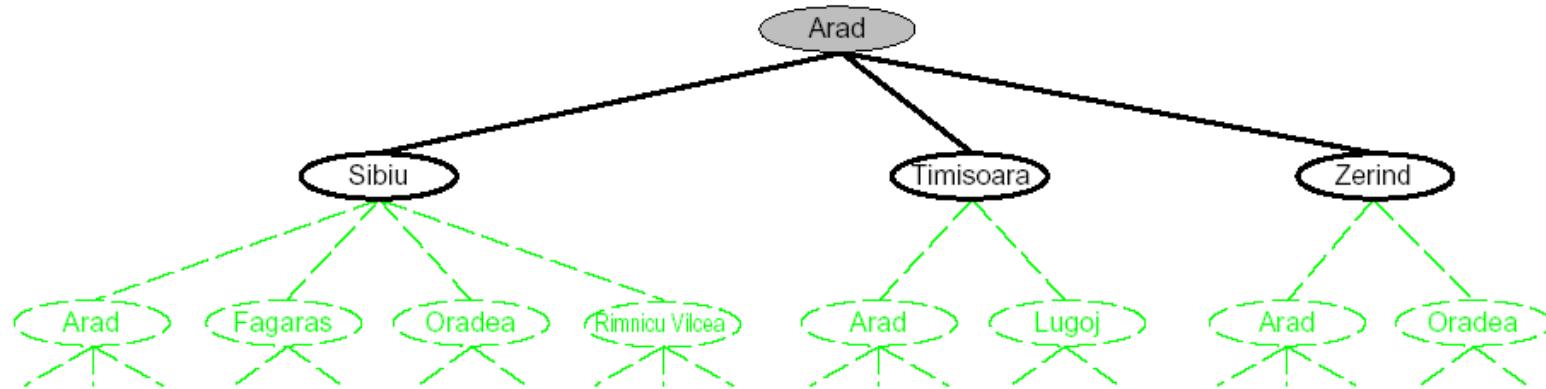


## Searching: *Tree Search*



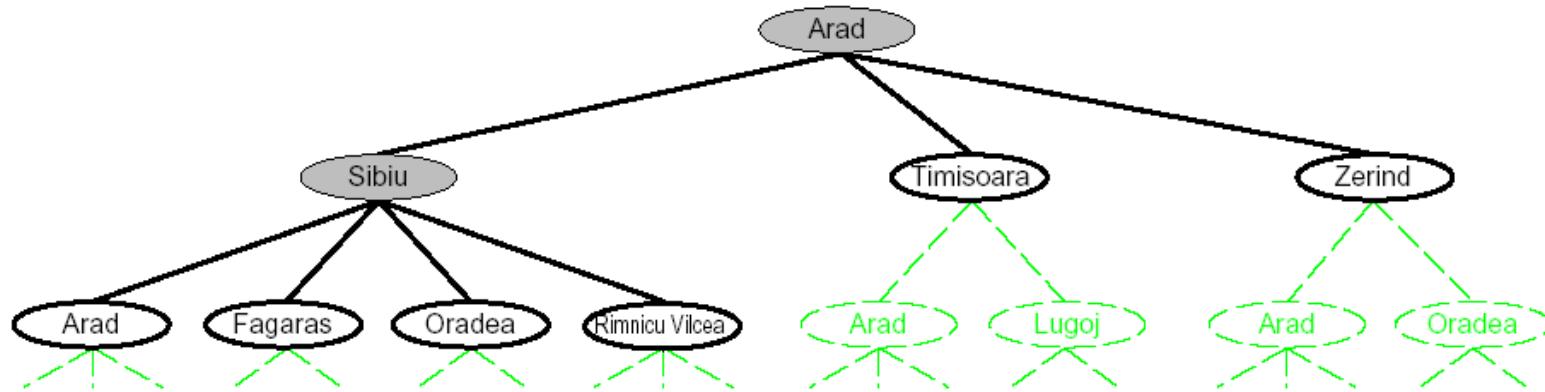


## Searching: *Tree Search*





## Searching: *Tree Search*



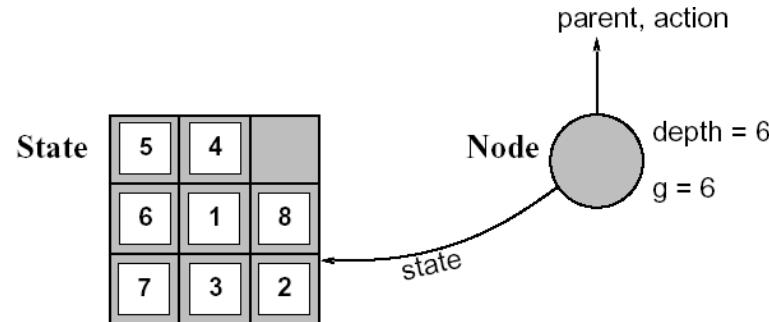


## Implementation: States as Nodes

A *state* is a (representation of) a physical configuration

A *node* is a data structure constituting part of a search tree    includes  
*parent, children, depth, path cost g(x)*

*States* do not have  
parents, children,  
depth, or path cost!



The **Expand** function creates new nodes, filling in the various fields and using the **SuccessorFn** of the problem to create the corresponding states.

# Implementation: General Tree Search

```
function TREE-SEARCH( problem, fringe ) returns a solution, or failure
  fringe  $\leftarrow$  INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node  $\leftarrow$  REMOVE-FRONT(fringe)
    if GOAL-TEST[problem] applied to STATE(node) succeeds return node
    fringe  $\leftarrow$  INSERTALL(EXPAND(node, problem), fringe)

function EXPAND( node, problem ) returns a set of nodes
  successors  $\leftarrow$  the empty set
  for each action, result in SUCCESSOR-FN[problem](STATE[node]) do
    s  $\leftarrow$  a new NODE
    PARENT-NODE[s]  $\leftarrow$  node; ACTION[s]  $\leftarrow$  action; STATE[s]  $\leftarrow$  result
    PATH-COST[s]  $\leftarrow$  PATH-COST[node] + STEP-COST(node, action, s)
    DEPTH[s]  $\leftarrow$  DEPTH[node] + 1
    add s to successors
  return successors
```



# Search Strategies



- A strategy is defined by picking the *order of node expansion*
- Strategies are evaluated along the following dimensions:
  - **Completeness** - does it always find a solution if one exists?
  - **Complexity** -number of nodes generated-expanded
  - **Space complexity** -maximum number of nodes in memory
  - **Optimality** - does it always find a least-cost solution?
- Time and space complexity are measured in terms of
  - **b**: maximum branching factor of the search tree
  - **d**: depth of the least-cost solution
  - **m**: maximum depth of the state space (may be  $\infty$ )



# Uninformed Search Strategies

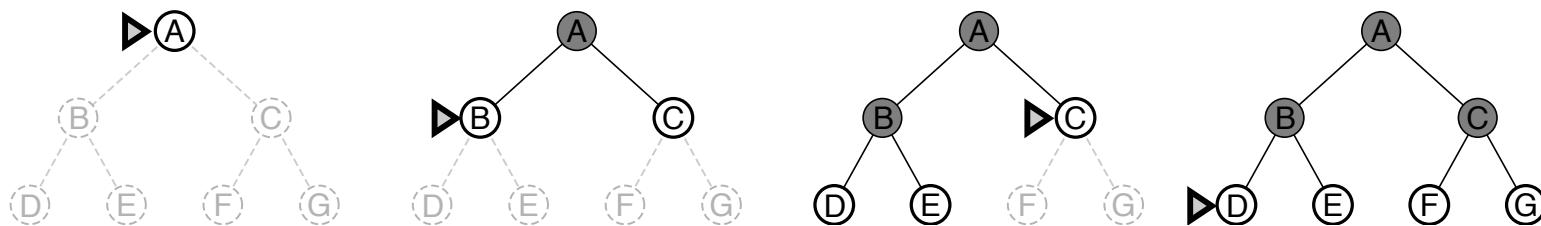
- *Uninformed* strategies use only the information available in the problem definition
  - Breadth-first search
  - Uniform-cost search
  - Depth-first search
  - Depth-limited search
  - Iterative deepening search



## Breadth-first search



- Expand the *shallowest* expanded Node
- Implementation:
  - **fringe** is a FIFO Queue, i.e., new successors go to the end





## Breadth-first search: Properties

### **Complete?**

- Yes (if  $b$  is finite)

### **Time?**

- $1+b+b^2+b^3+\dots+b^d + b(b^d-1) = O(b^{d+1})$ , i.e., exp. in  $d$

### **Space?**

- $O(b^{d+1})$  - (keeps every node in memory)

### **Optimal?**

- Yes (if cost = 1 per step); not optimal in general

Space is the big problem; can easily generate nodes at 10MB/sec, 24hrs = 860GB.



## Uniform-cost search



Expand least-cost unexpanded node

### ***Implementation:***

- `fringe` = queue ordered by path cost

Equivalent to breadth-first if step costs all equal

### ***Complete?***

- Yes, if step cost  $\geq \varepsilon$

### ***Time?***

- # of nodes with  $g \leq$  cost of optimal solution,  $O(b^{\lceil C^*/\varepsilon \rceil})$   
where  $C^*$  is the cost of the optimal solution

### ***Space?***

- # of nodes with  $g \leq$  cost of optimal solution,  $O(b^{\lceil C^*/\varepsilon \rceil})$

### ***Optimal?***

- Yes, nodes expanded in increasing order of  $g(n)$



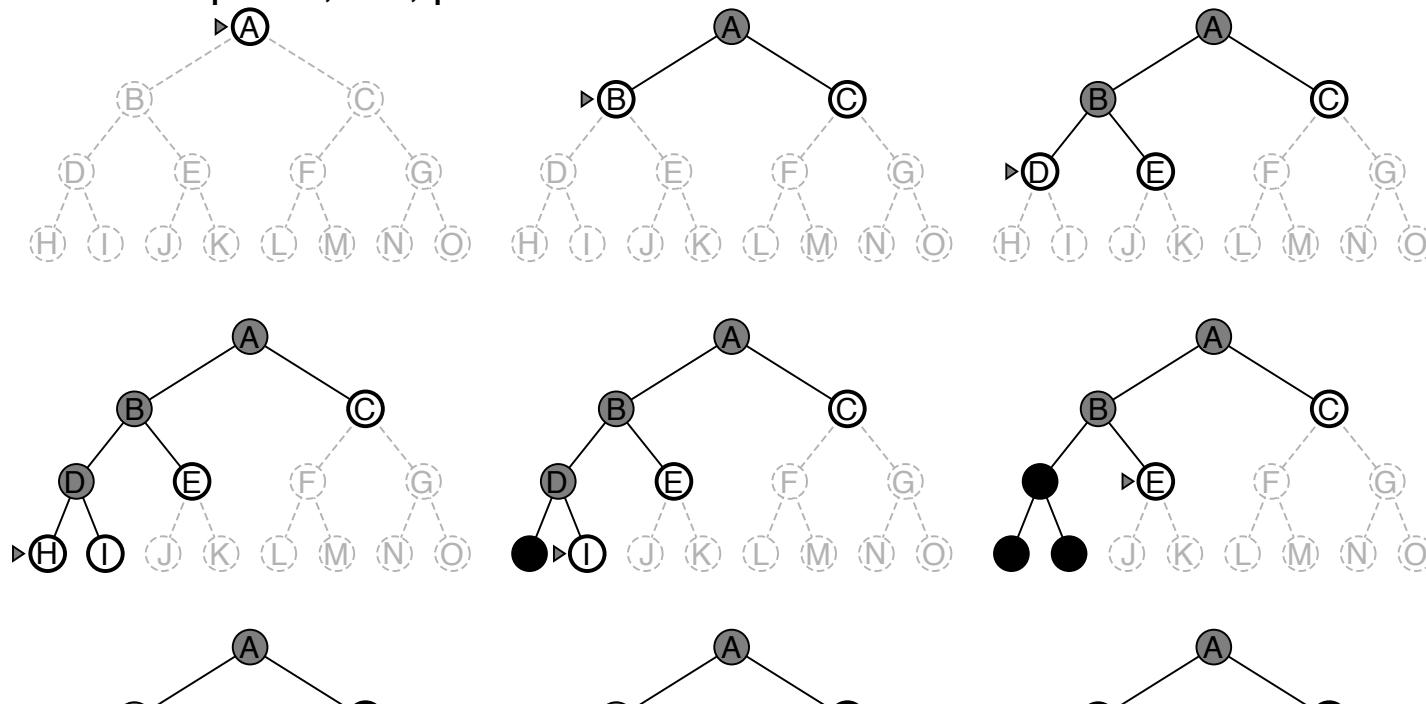
## Depth-first search



Expand deepest unexpanded node

***Implementation:***

**fringe** = LIFO queue, i.e., put successors at front





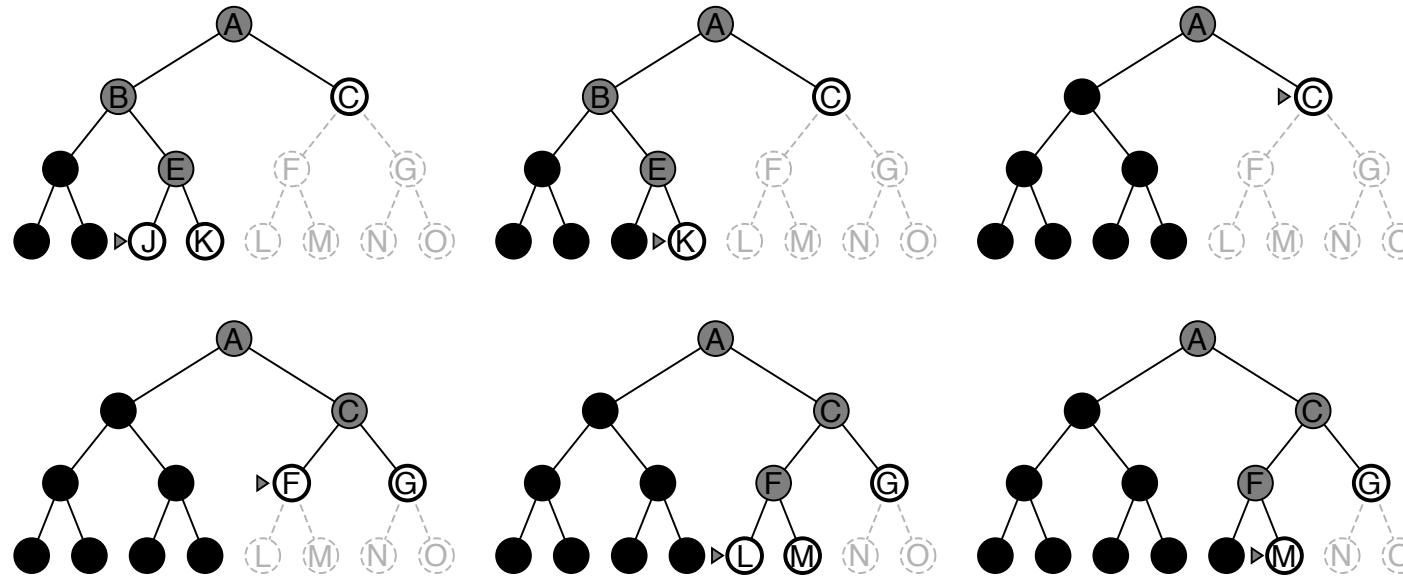
## Depth-first search (cont.)



Expand deepest unexpanded node

***Implementation:***

**fringe** = LIFO queue, i.e., put successors at front





# Depth-first search: Properties

## ***Complete?***

- No: fails in infinite-depth spaces, spaces with loops
- Modify to avoid repeated states along path
- → complete in finite spaces

## ***Time?***

- $O(b^m)$ : terrible if **m** is much larger than **d**  
but if solutions are dense, may be much faster than breadth-first

## ***Space?***

- $O(bm)$ , i.e., linear space!

## ***Optimal?***

- No



## Depth-limited search

= depth-first search with depth limit  $l$ ,  
i.e., nodes at depth  $l$  have no successors

```
function DEPTH-LIMITED-SEARCH( problem, limit) returns soln/fail/cutoff
    RECURSIVE-DLS(MAKE-NODE(INITIAL-STATE[problem]), problem, limit)
    
function RECURSIVE-DLS(node, problem, limit) returns soln/fail/cutoff
    cutoff-occurred?  $\leftarrow$  false
    if GOAL-TEST[problem](STATE[node]) then return node
    else if DEPTH[node] = limit then return cutoff
    else for each successor in EXPAND(node, problem) do
        result  $\leftarrow$  RECURSIVE-DLS(successor, problem, limit)
        if result = cutoff then cutoff-occurred?  $\leftarrow$  true
        else if result  $\neq$  failure then return result
    if cutoff-occurred? then return cutoff else return failure
```



## Iterative deepening search

```
function ITERATIVE-DEEPENING-SEARCH(problem) returns a solution
    inputs: problem, a problem

    for depth  $\leftarrow 0$  to  $\infty$  do
        result  $\leftarrow$  DEPTH-LIMITED-SEARCH(problem, depth)
        if result  $\neq$  cutoff then return result
    end
```



## Iterative deepening search $l = 0$

L

Limit = 0



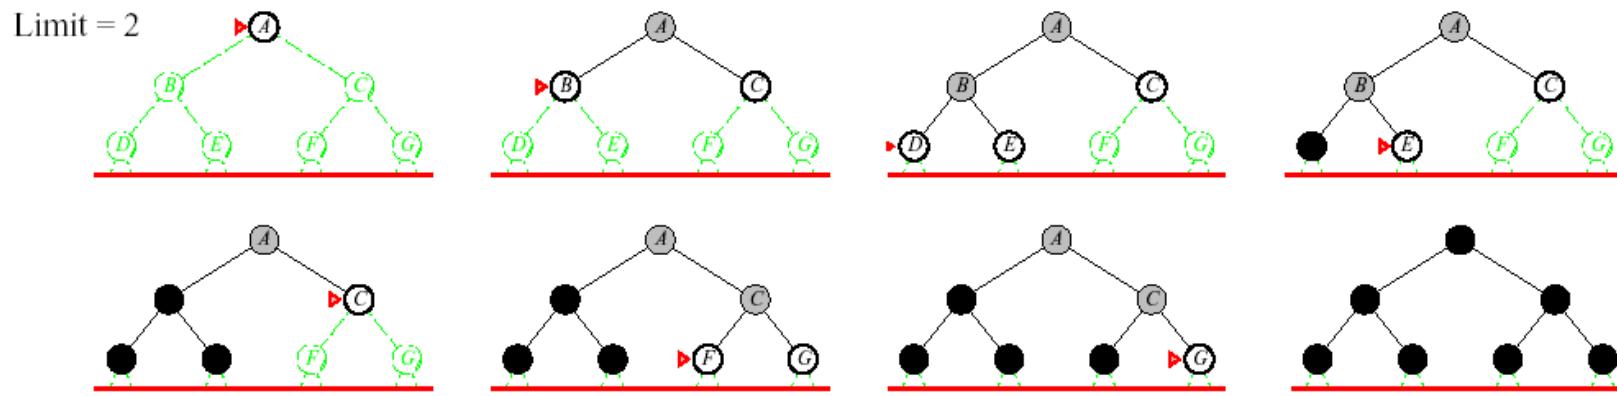


## Iterative deepening search $l = 1$

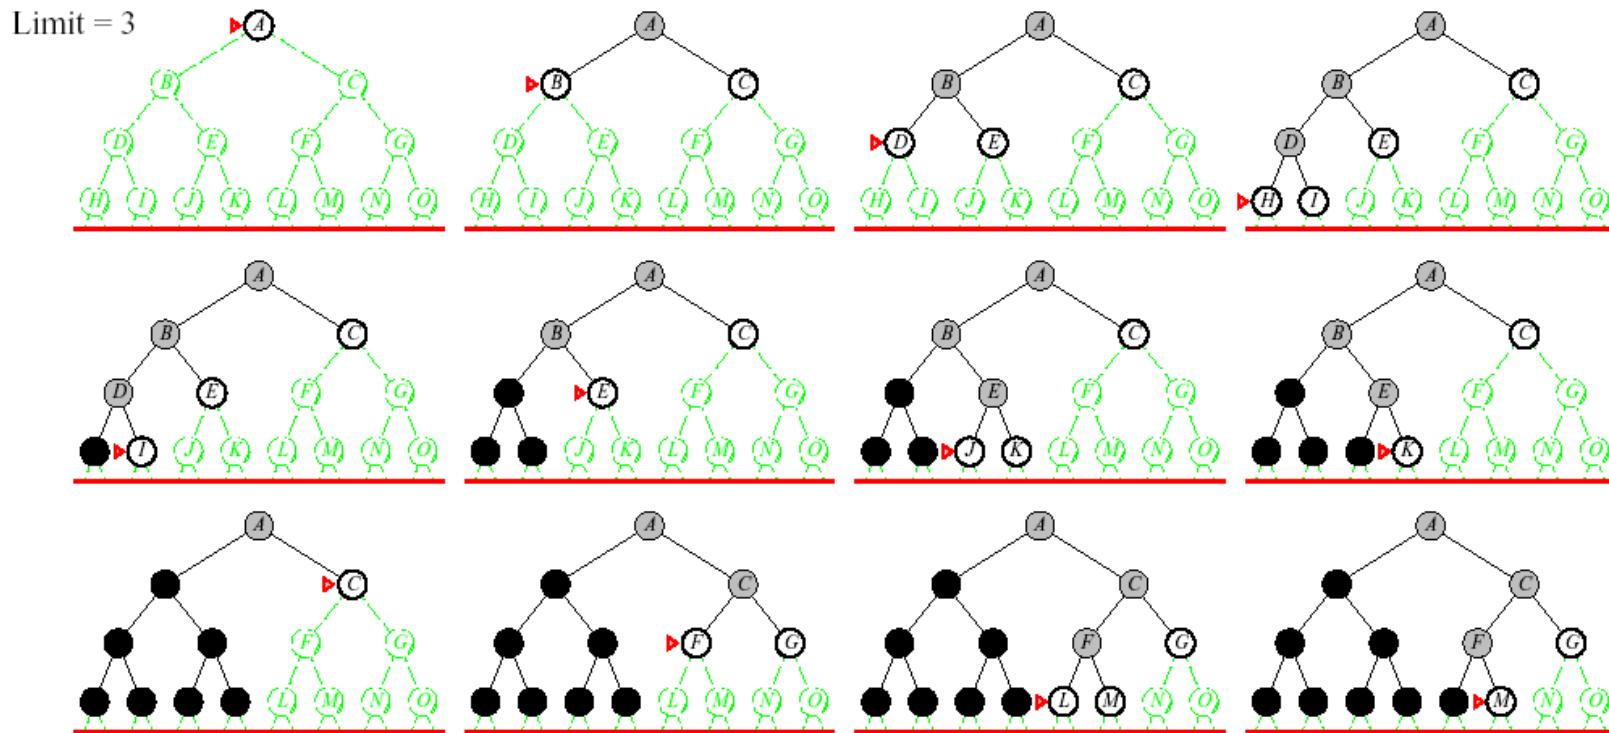
Limit = 1



# Iterative deepening search $l = 2$



## Iterative deepening search $l = 3$





## Iterative deepening search: Properties



### ***Complete?***

- Yes

### ***Time?***

- $(d+1)b^0 + d b^1 + (d-1)b^2 + \dots + b^d = O(b^d)$

### ***Space?***

- $O(bd)$

### ***Optimal?***

- Yes, if step cost = 1  
Can be modified to explore uniform-cost tree

Numerical comparison for  $b=10$  and  $d=5$ , solution at far right:

$$N(IDS) = 50 + 400 + 3,000 + 20,000 + 100,000 = 123,450$$

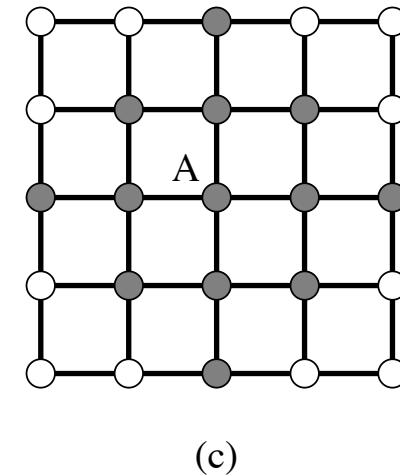
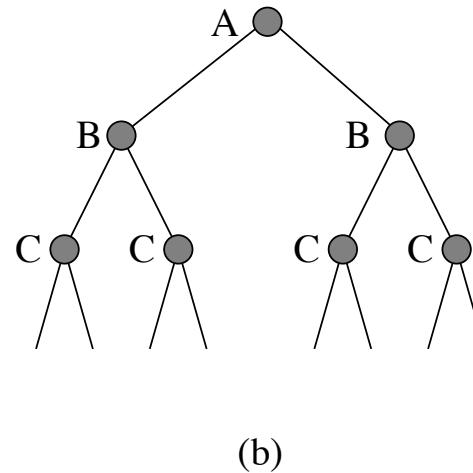
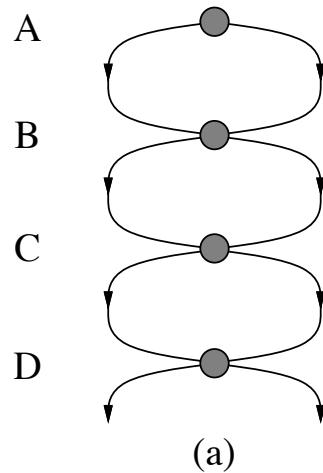
$$N(BFS) = 10 + 100 + 1,000 + 10,000 + 100,000 + 999,990 = 1,111,100$$

# Comparison of Algorithms

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening
Complete?	Yes*	Yes*	No	Yes, if $l \geq d$	Yes
Time	$b^{d+1}$	$b^{\lceil C^*/\epsilon \rceil}$	$b^m$	$b^l$	$b^d$
Space	$b^{d+1}$	$b^{\lceil C^*/\epsilon \rceil}$	$bm$	$bl$	$bd$
Optimal?	Yes*	Yes*	No	No	Yes*

# Repeated States

- Failure to detect repeated states can turn a linear problem into an exponential one!





# Graph Search

```
function GRAPH-SEARCH( problem, fringe) returns a solution, or failure
    closed  $\leftarrow$  an empty set
    fringe  $\leftarrow$  INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
    loop do
        if fringe is empty then return failure
        node  $\leftarrow$  REMOVE-FRONT(fringe)
        if GOAL-TEST[problem](STATE[node]) then return node
        if STATE[node] is not in closed then
            add STATE[node] to closed
            fringe  $\leftarrow$  INSERTALL(EXPAND(node, problem), fringe)
    end
```



## Summary



- Problem formulation usually requires abstracting away real-world details to define a state space that can feasibly be explored
- Variety of uninformed search strategies
- Iterative deepening search uses only linear space and not much more time than other uninformed algorithms



## Overview – Part 1



- Problem Definition?
- Search
  - Breadth first, Depth first, Iterative deepening
- **Informed (or Intelligent) Search**
  - Best First
  - A\*, Hill climbing, Simulated annealing
- Constraint Satisfaction
  - Forward checking, Constraint propagation
- Adversarial Search (Games)



## Informed Search Strategies



- *Informed* strategies use some information about the problem (a heuristic) to improve performance
  - Best-first search
  - A\* search
  - Heuristics
  - Hill-climbing
  - Simulated annealing



# Review: Tree Search

---

```
function TREE-SEARCH( problem, fringe ) returns a solution, or failure
  fringe  $\leftarrow$  INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node  $\leftarrow$  REMOVE-FRONT(fringe)
    if GOAL-TEST[problem] applied to STATE(node) succeeds return node
    fringe  $\leftarrow$  INSERTALL(EXPAND(node, problem), fringe)
```

---

```
function EXPAND( node, problem ) returns a set of nodes
  successors  $\leftarrow$  the empty set
  for each action, result in SUCCESSOR-FN[problem](STATE[node]) do
    s  $\leftarrow$  a new NODE
    PARENT-NODE[s]  $\leftarrow$  node; ACTION[s]  $\leftarrow$  action; STATE[s]  $\leftarrow$  result
    PATH-COST[s]  $\leftarrow$  PATH-COST[node] + STEP-COST(node, action, s)
    DEPTH[s]  $\leftarrow$  DEPTH[node] + 1
    add s to successors
  return successors
```

- A strategy is defined by picking the order of node expansion
- How can we choose a “smart order”?



## Best First Search



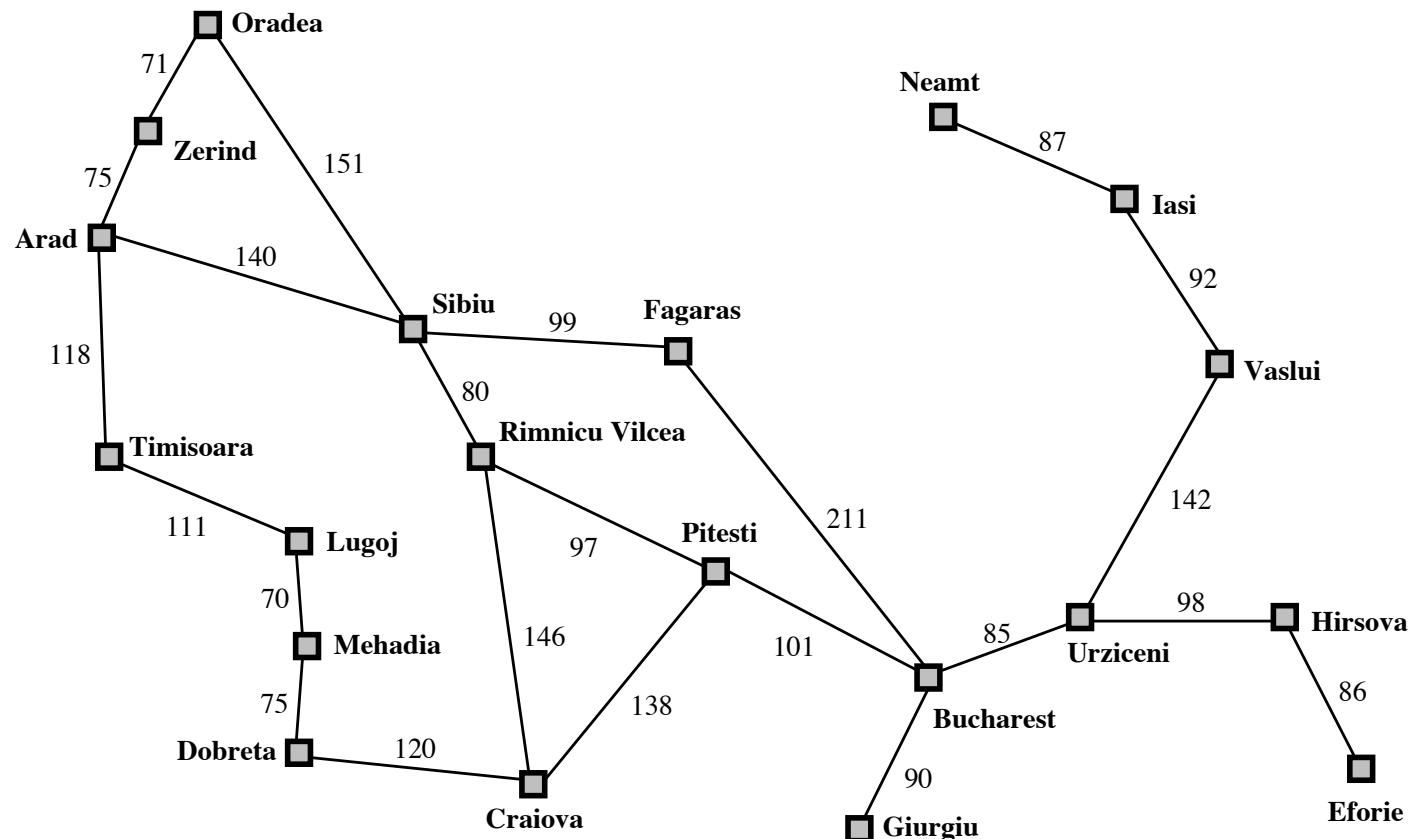
### Idea:

- use an **evaluation** function for each node as an estimate of “**desirability**”  
→ Expand most desirable unexpanded node

### Implementation:

- **fringe** is a queue sorted in decreasing order of desirability
- Special cases:
  - greedy search
  - A\* search

# Romania with Step Cost in km



	Straight-line distance to Bucharest
<b>Arad</b>	366
<b>Bucharest</b>	0
<b>Craiova</b>	160
<b>Dobreta</b>	242
<b>Eforie</b>	161
<b>Fagaras</b>	178
<b>Giurgiu</b>	77
<b>Hirsova</b>	151
<b>Iasi</b>	226
<b>Lugoj</b>	244
<b>Mehadia</b>	241
<b>Neamt</b>	234
<b>Oradea</b>	380
<b>Pitesti</b>	98
<b>Rimnicu Vilcea</b>	193
<b>Sibiu</b>	253
<b>Timisoara</b>	329
<b>Urziceni</b>	80
<b>Vaslui</b>	199
<b>Zerind</b>	374



## Greedy search



- Evaluation function  $h(n)$  (heuristic) = estimate of cost from  $n$  to the closest goal
- E.g.,  $h_{SLD}(n)$  = straight-line distance from  $n$  to Bucharest
- Greedy search expands the node that *appears* to be closest to goal

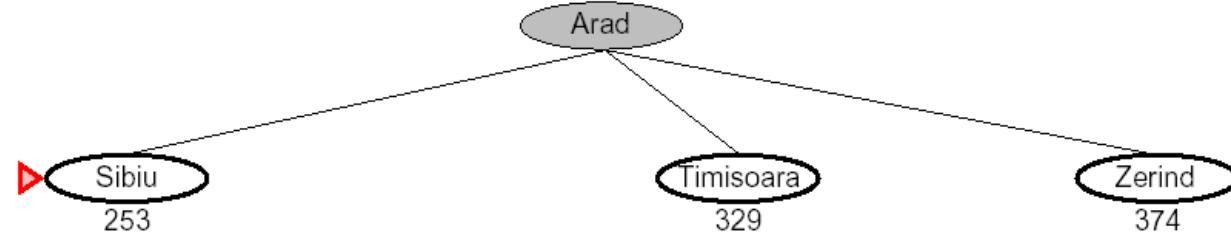


# Greedy Search



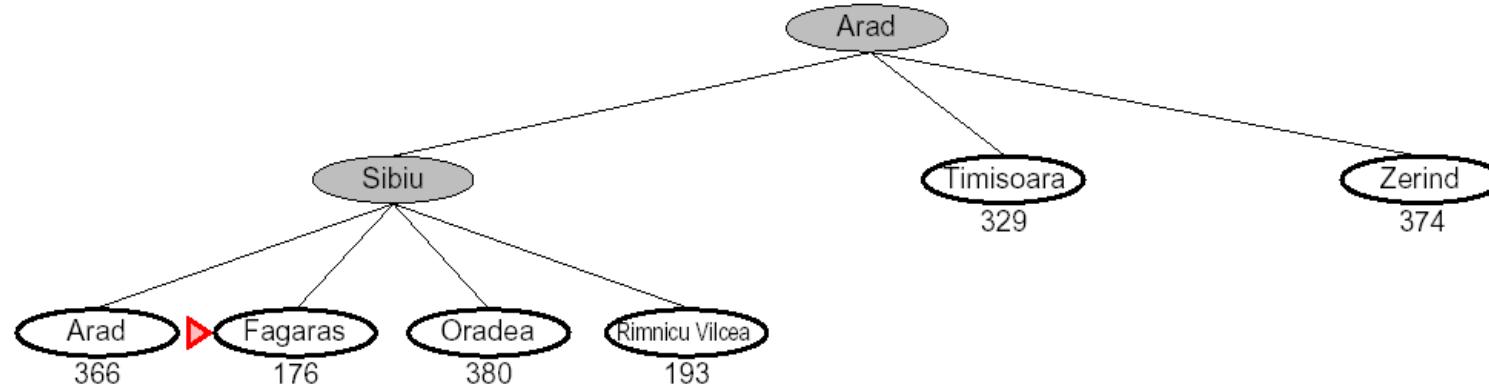


## Greedy Search



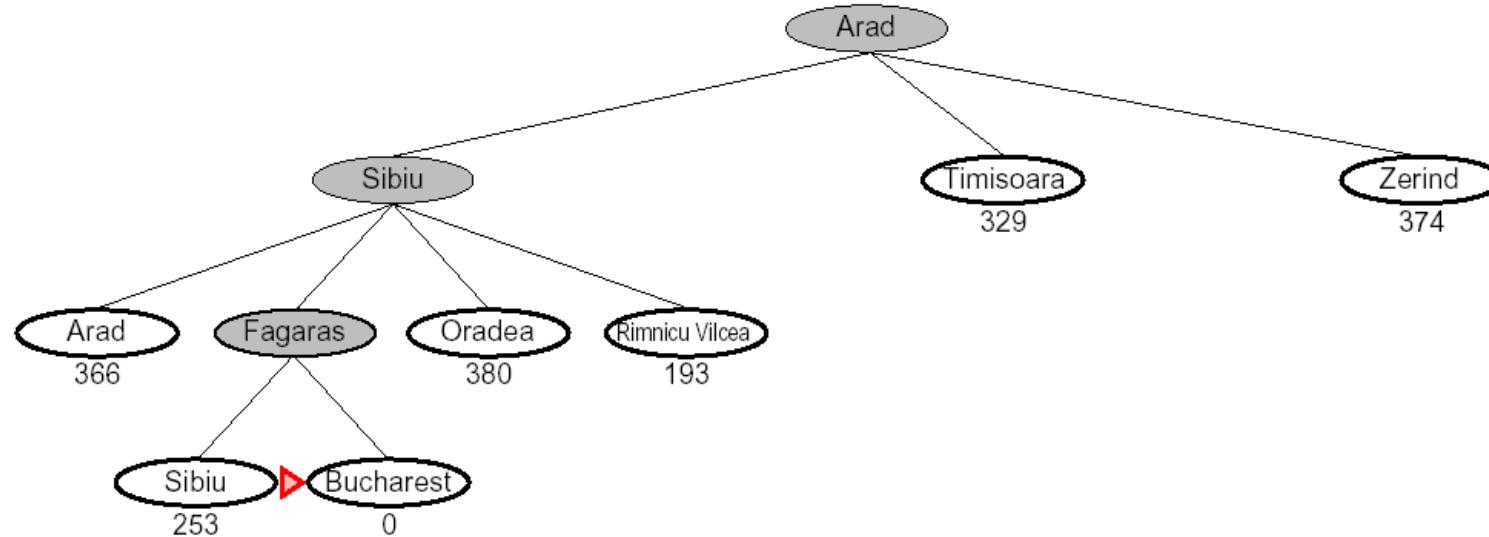


# Greedy Search





# Greedy Search





## Greedy Search: Properties



### ***Complete?***

- No – can get stuck in loops,  
e.g., lași → Neamț → lași → Neamț →
- Complete in finite space with repeated-state checking

### ***Time?***

- $O(b^m)$ , but a good heuristic can give dramatic improvement

### ***Space?***

- $O(b^m)$ , keeps all nodes in memory

### ***Optimal***

- No



## A\* Search



**Idea:** avoid expanding paths that are already expensive

- Evaluation function  $f(n) = g(n) + h(n)$
- $g(n)$  = cost so far to reach  $n$
- $h(n)$  = estimated cost to goal from  $n$
- $f(n)$  = estimated total cost of path through  $n$  to goal
- A\* search uses an *admissible* heuristic  
i.e.,  $h(n) \leq h^*(n)$ , where  $h^*(n)$  is the *true* cost from  $n$   
(Also require  $h(n) \geq 0$ , so  $h(G) = 0$  for any goal  $G$ )
- E.g.,  $h_{SLD}(n)$  never overestimates the actual road distance

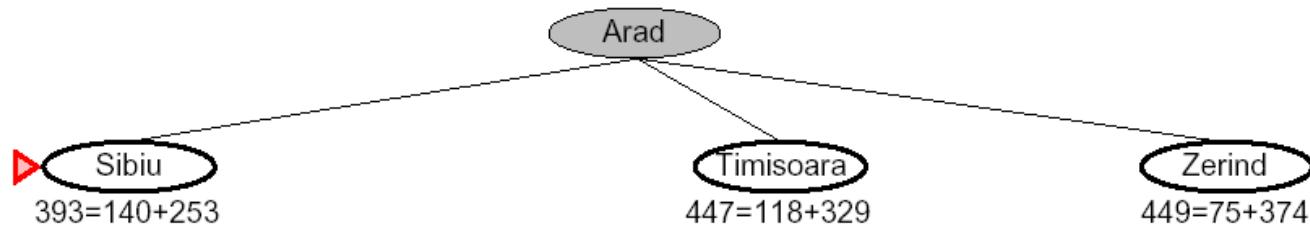
**Theorem:** A\* search is optimal

# A\* Search

► Arad  
 $366=0+366$

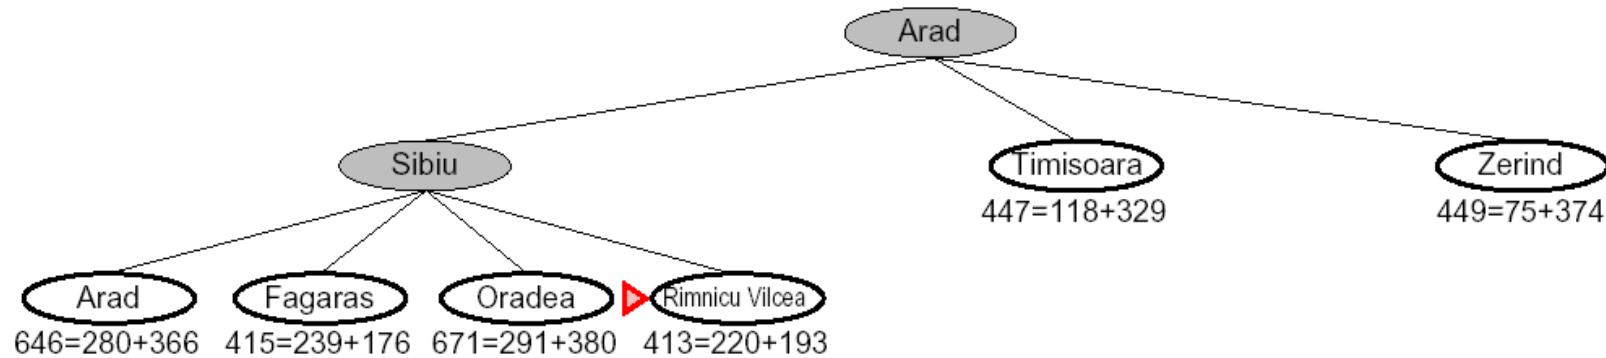


## A\* Search



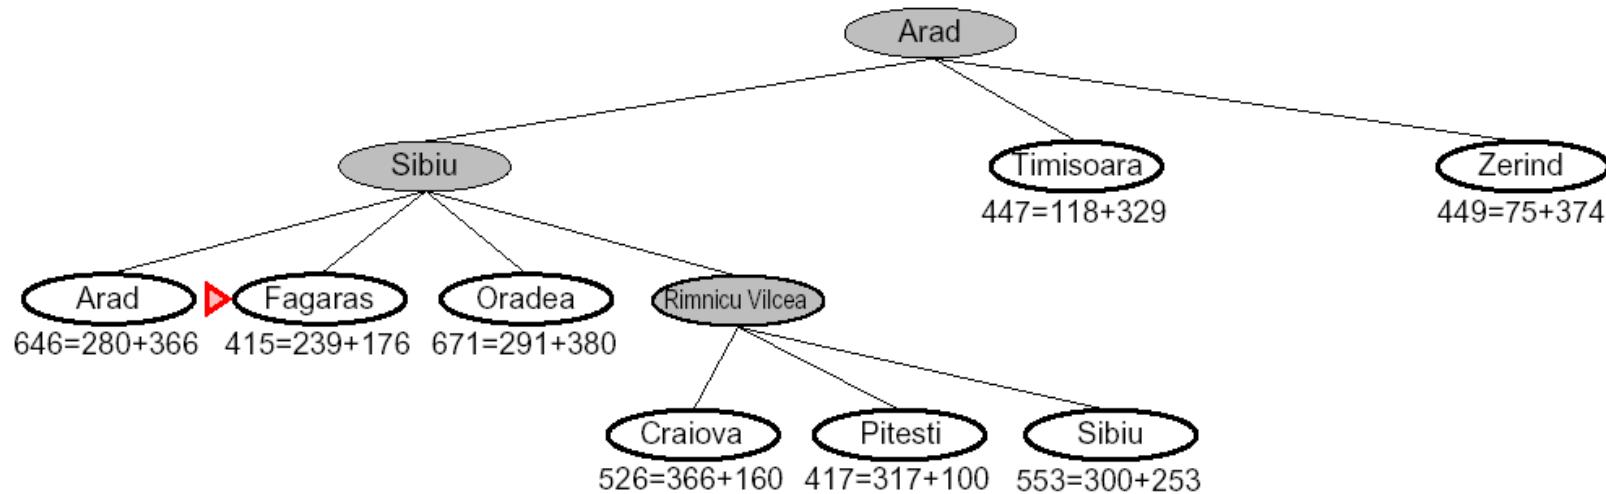


## A\* Search



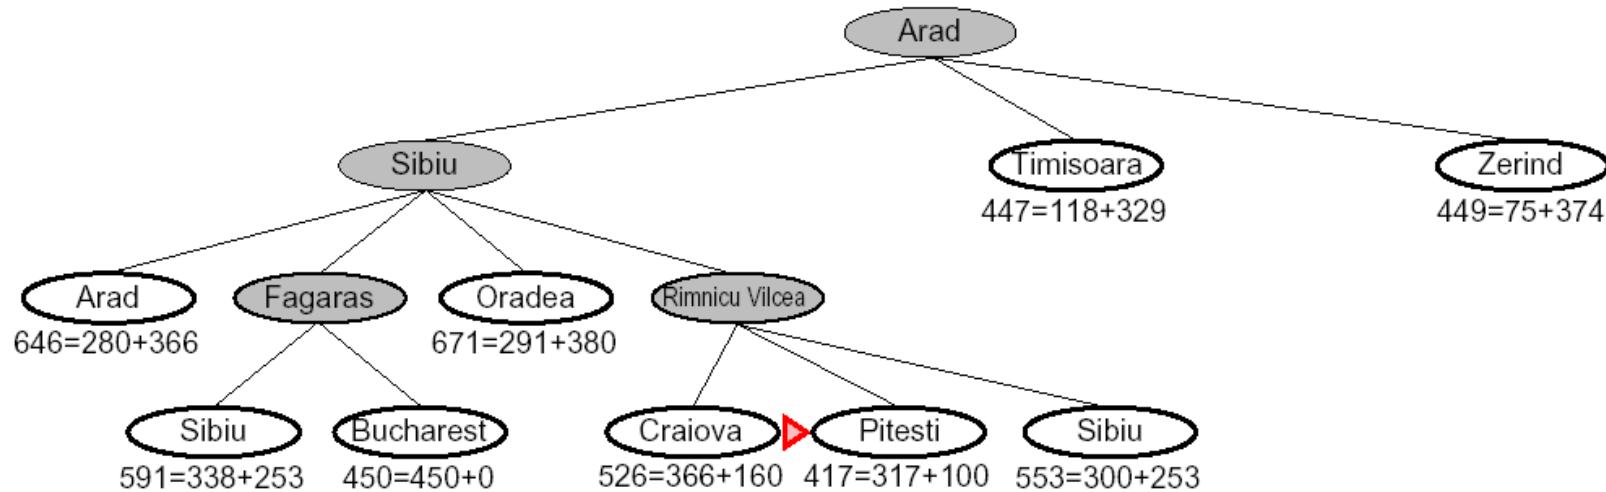


## A\* Search



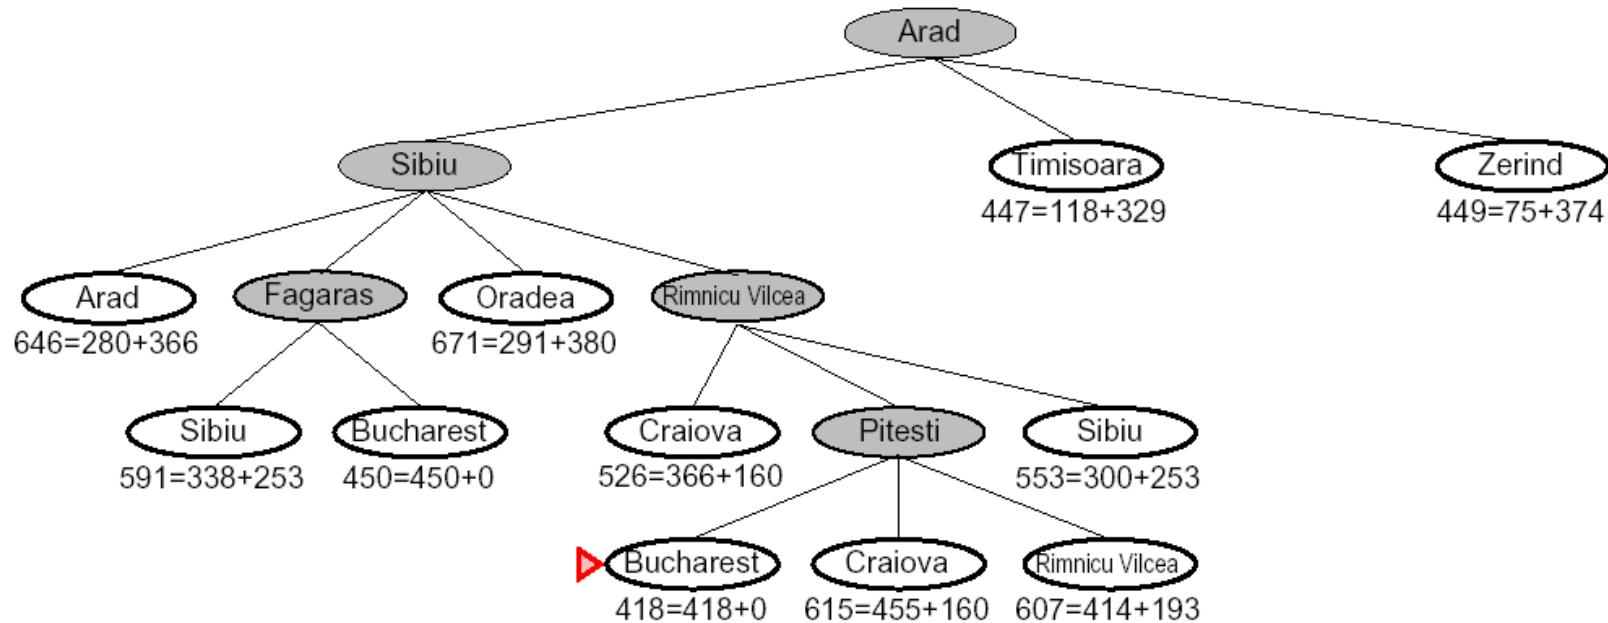


# A\* Search





# A\* Search

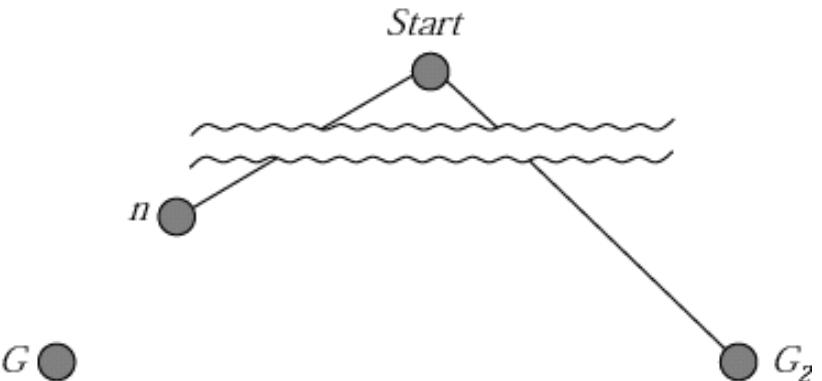




## Optimality of A\* (standard proof)

- Suppose some suboptimal goal  $G_2$  has been generated and is in the queue. Let  $n$  be an unexpanded node on a shortest path to an optimal goal  $G_1$ .

$$\begin{aligned}f(G_2) &= g(G_2) \text{ since } h(G_2) = 0 \\&> g(G_1) \text{ since } G_2 \text{ is suboptimal} \\&\geq f(n) \text{ since } h \text{ is admissible}\end{aligned}$$

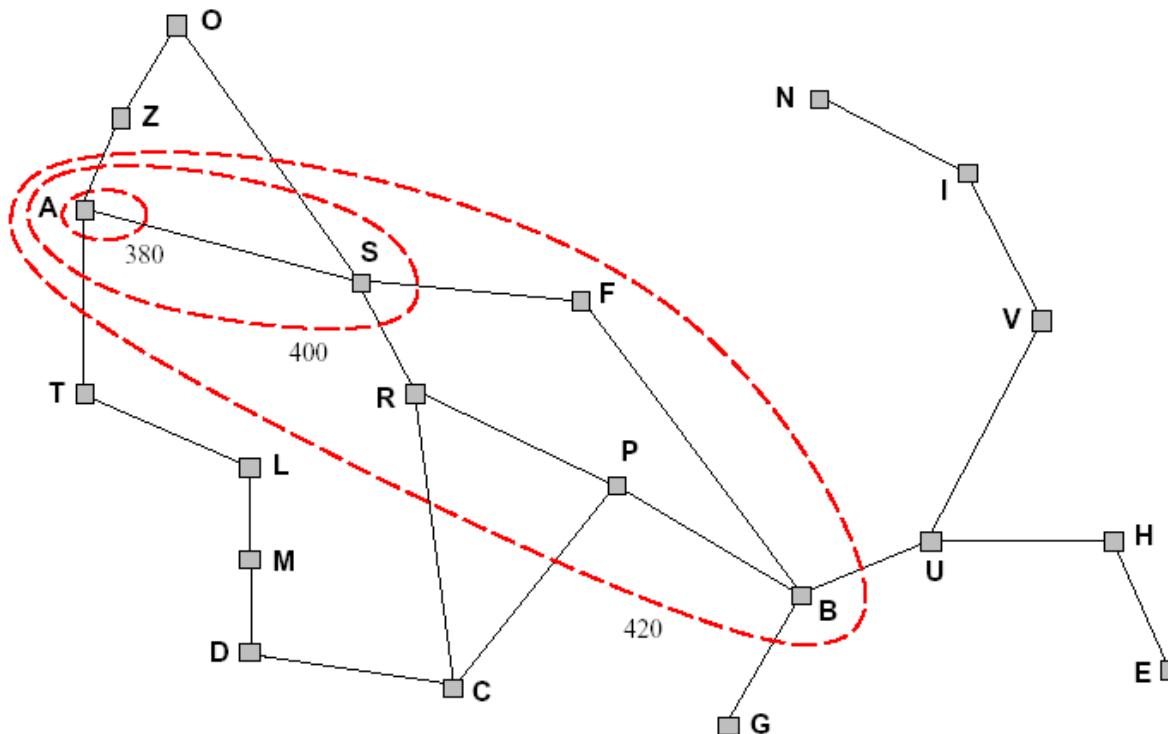


- Since  $f(G_2) > f(n)$ ,  $A^*$  will never select  $G_2$  for expansion

└

## Optimality of A\* (more useful)

- *Lemma:* A\* expands nodes in order of increasing  $f$  value\*
- Gradually adds “ $f$ -contours” of nodes (cf. breadth-first adds layers)
- Contour  $i$  has all nodes with  $f = f_i$ , where  $f_i < f_{i+1}$





## A\* search: Properties



### **Complete?**

- Yes,  
unless there are infinitely many nodes with  $f \leq f(G)$

### **Time?**

- Exponential in [relative error in  $h$  \* length of soln.]

### **Space?**

- Keeps all nodes in memory

### **Optimal?**

- Yes, cannot expand  $f_{i+1}$  until  $f_i$  is finished
- A\* expands all nodes with  $f(n) < c^*$
- A\* expands some nodes with  $f(n) = c^*$
- A\* expands no nodes with  $f(n) > c^*$



## Admissible Heuristics



E.g., for the 8-puzzle:

- $h_1(n)$  = number of misplaced tiles
- $h_2(n)$  = total *Manhattan* distance (i.e., no. of squares from desired location of each tile)

7	2	4
5		6
8	3	1

Start State

1	2	3
4	5	6
7	8	

Goal State

- $h_1(S) = 7$
- $h_2(S) = 4+0+3+3+1+0+2+1 = 14$



## Dominance



- If  $h_2(n) \geq h_1(n)$  for all  $n$  (both admissible) then  $h_2(n)$  **dominates**  $h_1(n)$  and is better for search

Typical search costs:

$d=14$       IDS      = 3,473,941 nodes

$A^*(h_1)$       = 539 nodes

$A^*(h_2)$       = 113 nodes

$d=24$       IDS      approx 54,000,000,000 nodes

$A^*(h_1)$       = 39,135 nodes

$A^*(h_2)$       = 1,641 nodes



## Relaxed Problems



- Admissible heuristics can be derived from the *exact* solution cost of a *relaxed* version of the problem
- If the rules of the 8-puzzle are relaxed so that a tile can move *anywhere*, then  $h_1(n)$  gives the shortest solution
- If the rules are relaxed so that a tile can move to *any adjacent square*, then  $h_2(n)$  gives the shortest solution
- Key point: the optimal solution cost of a relaxed problem is no greater than the optimal solution cost of the real problem



# Iterative Improvement Algorithms



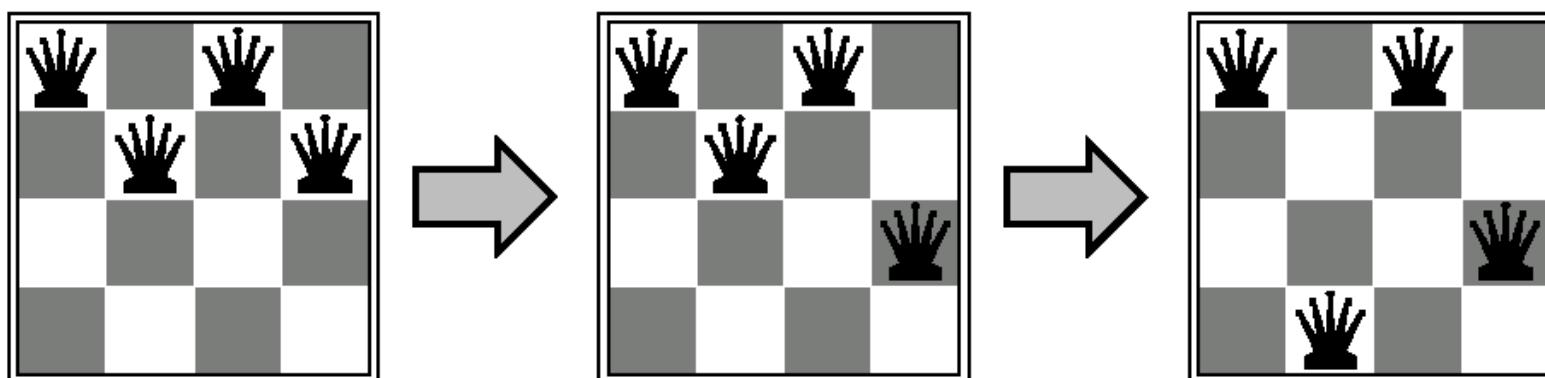
- In many optimization problems, *path* is irrelevant; the goal state itself is the solution
- Then state space = set of “complete” configurations;
  - find *optimal* configuration, e.g., TSP
  - or, find configuration satisfying constraints, e.g., timetable
- In such cases, can use *iterative improvement* algorithms;
- keep a single “current” state, try to improve it
- Constant space, **suitable** for online as well as offline search



## Example: n-queens



- Put  $n$  queens on an  $n * n$  board with no two queens on the same row, column, or diagonal
- Move a queen to reduce number of conflicts





## Hill-climbing (or gradient ascent/descent)

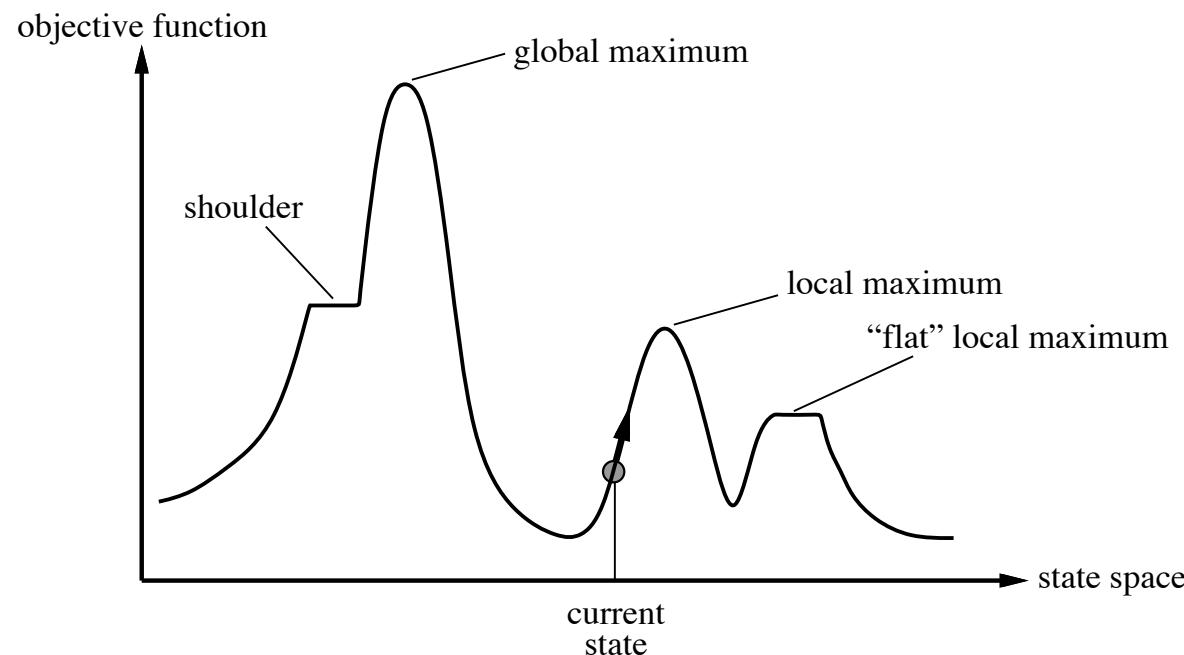
- “Like climbing Everest in thick fog with amnesia”

```
function HILL-CLIMBING(problem) returns a state that is a local maximum
  inputs: problem, a problem
  local variables: current, a node
                  neighbor, a node
  current  $\leftarrow$  MAKE-NODE(INITIAL-STATE[problem])
  loop do
    neighbor  $\leftarrow$  a highest-valued successor of current
    if VALUE[neighbor] < VALUE[current] then return STATE[current]
    current  $\leftarrow$  neighbor
  end
```

# └ Hill-climbing (cont.)

## └ (or gradient ascent/descent)

**Problem:** depending on initial state, can get stuck on local maxima





## Simulated annealing



**Idea:** escape local maxima by allowing some “bad” moves *but gradually decrease their size and frequency*

**function** SIMULATED-ANNEALING(*problem, schedule*) **returns** a solution state

**inputs:** *problem*, a problem

*schedule*, a mapping from time to “temperature”

**local variables:** *current*, a node

*next*, a node

*T*, a “temperature” controlling prob. of downward steps

*current*  $\leftarrow$  MAKE-NODE(INITIAL-STATE[*problem*])

**for** *t*  $\leftarrow$  1 **to**  $\infty$  **do**

*T*  $\leftarrow$  *schedule*[*t*]

**if** *T* = 0 **then return** *current*

*next*  $\leftarrow$  a randomly selected successor of *current*

$\Delta E \leftarrow$  VALUE[*next*] – VALUE[*current*]

**if**  $\Delta E > 0$  **then** *current*  $\leftarrow$  *next*

**else** *current*  $\leftarrow$  *next* only with probability  $e^{\Delta E/T}$



## Simulated annealing: Properties

- At fixed “temperature” T, state occupation probability reaches Boltzman distribution

$$p(x) = \alpha^{\frac{E(x)}{kT}}$$

- T decreased slowly enough → always reach best state
- Is this necessarily an interesting guarantee?
- Devised by Metropolis et al., 1953, for physical process modelling
- Widely used in VLSI layout, airline scheduling, etc.

# Practical AI (a.k.a. Business Intelligence)

## ***Part 1- Intelligent Search***

Search

(Intelligent) Search

**Constraint Satisfaction**

Adversarial Search (Games)



**Universität  
Zürich<sup>UZH</sup>**



Dynamic and Distributed  
Information Systems





## Overview – Part 1



- Problem Definition?
- Search
  - Breadth first, Depth first, Iterative deepening
  - Best First
- Informed (or Intelligent) Search
  - A\*, Hill climbing, Simulated annealing
- **Constraint Satisfaction**
  - Forward checking, Constraint propagation
- Adversarial Search (Games)



## CSP: Outline



- CSP examples
- Backtracking search for CSPs
- Problem structure and problem decomposition
- Local search for CSPs

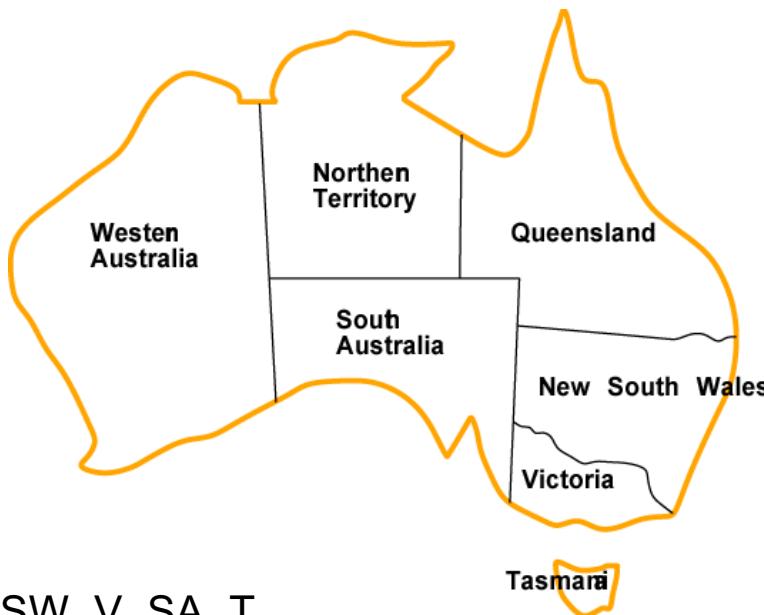


# Constraint satisfaction problems (CSPs)

- Standard search problem:
  - **state** is a “black box” – any old data structure that supports goal test, eval, successor
- CSP:
  - **state** is defined by *variables*  $X_i$  with *values* from domain  $D_i$
  - **goal test** is a set of *constraints* specifying allowable combinations of values for subsets of variables
- Simple example of a *formal representation language*
- Allows useful *general-purpose* algorithms with more power than standard search algorithms



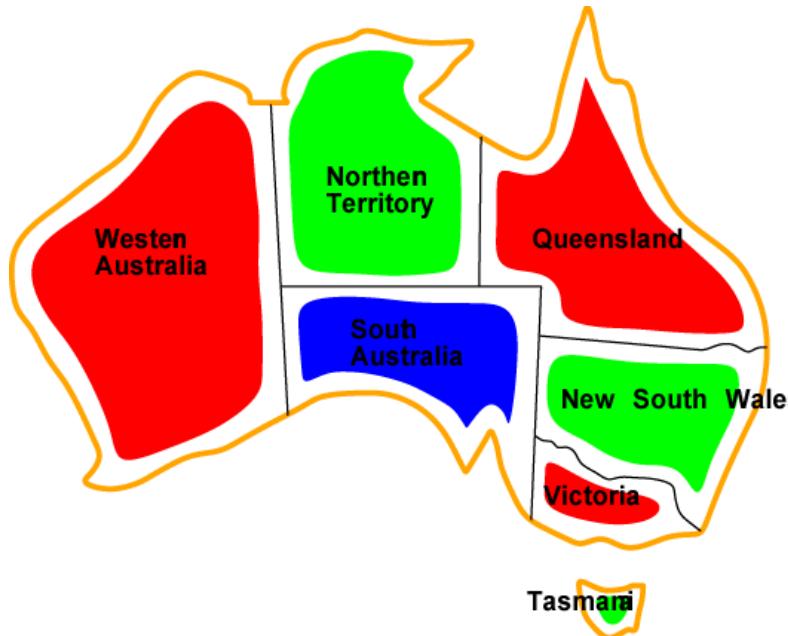
## Example: Map-Coloring



- **Variables:** WA, NT, Q, NSW, V, SA, T
- **Domains:**  $D_i = \{\text{red, green, blue}\}$
- **Constraints:** adjacent regions must have different colors
  - e.g.,  $WA \neq NT$  (if the language allows this), or  
 $(WA,NT) \in \{(\text{red, green}), (\text{red, blue}), (\text{green, red}), (\text{green, blue}), \dots\}$



## Example: Map-Coloring (contd.)



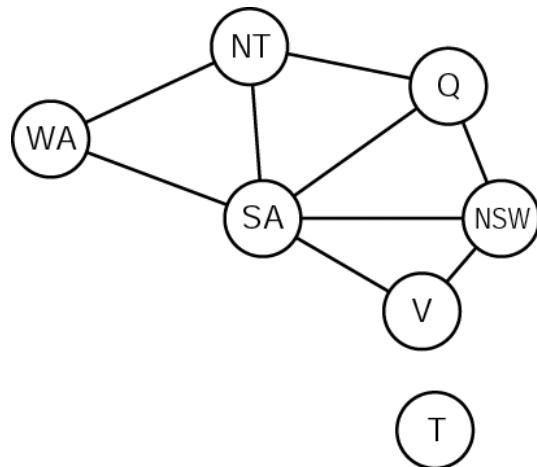
- **Solutions** are assignments satisfying all constraints,
  - e.g., {WA = red, NT = green, Q = red, NSW = green, V = red, SA = blue, T = green}



## Constraint graph



- *Binary CSP*: each constraint relates at most two variables
- *Constraint graph*: nodes are variables, arcs show constraints



- General-purpose CSP algorithms use the graph structure to speed up search.
  - E.g., Tasmania is an independent subproblem!



# Varieties of CSPs



- Discrete variables
  - finite domains; size  $d \Rightarrow O(d^n)$  complete assignments
    - e.g., Boolean CSPs, incl. Boolean satisfiability (NP-complete)
  - infinite domains (integers, strings, etc.)
    - e.g., job scheduling, variables are start/end days for each job
    - need a *constraint language*, e.g.,  $\text{StartJob}_1 + 5 \leq \text{StartJob}_3$
    - *linear* constraints solvable, *nonlinear* undecidable
- Continuous variables
  - e.g., start/end times for Hubble Telescope observations
  - linear constraints solvable in poly time by LP methods



## Varieties of constraints

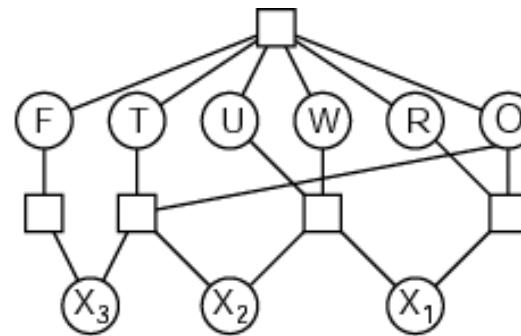


- *Unary* constraints involve a single variable,
    - e.g.,  $SA \neq green$
  - *Binary* constraints involve pairs of variables,
    - e.g.,  $SA \neq WA$
  - *Higher-order* constraints involve 3 or more variables,
    - e.g., cryptarithmetic column constraints
  - *Preferences* (soft constraints), e.g., *red* is better than *green* often representable by a cost for each variable assignment
- constrained optimization problems



## Example: Cryptarithmetic

$$\begin{array}{r} \text{T} \ \text{W} \ \text{O} \\ + \ \text{T} \ \text{W} \ \text{O} \\ \hline \text{F} \ \text{O} \ \text{U} \ \text{R} \end{array}$$



- *Variables:*  $F, T, U, W, R, O, X_1, X_2, X_3$
- *Domains:*  $\{0, 1, 2, 3, 4, 5, 6, 7, 8, 9\}$
- *Constraints:*
  - $\text{alldiff}(F, T, U, W, R, O)$
  - $O + O = R + 10 * X_1$ , etc.



## Real-world CSPs



- Assignment problems
    - e.g., who teaches what class
  - Timetabling problems
    - e.g., which class is offered when and where?
  - Hardware configuration
  - Spreadsheets
  - Transportation scheduling
  - Factory scheduling
  - Floorplanning
- 
- *Note:* that many real-world problems involve real-valued variables



# Standard search formulation (incremental)

Let's start with the straightforward, dumb approach, then fix it

States are defined by the values assigned so far

- **Initial state:** the empty assignment, {}
- **Successor function:** assign a value to an unassigned variable that does not conflict with current assignment.
- $\Rightarrow$  fail if no legal assignments (not fixable!)
- **Goal test:** the current assignment is complete

- 1) This is the same for all CSPs!
- 2) Every solution appears at depth  $n$  with  $n$  variables  
 $\Rightarrow$  use depth-first search
- 3) Path is irrelevant, so can also use complete-state formulation
- 4)  $b = (n-l)d$  at depth  $l$ , hence  $n!d^n$  leaves!!!!



## Backtracking search



- Variable assignments are *commutative*, i.e.,
  - [ $WA = \text{red}$  then  $NT = \text{green}$ ] same as [ $NT = \text{green}$  then  $WA = \text{red}$ ]
- Only need to consider assignments to a single variable at each node
  - ⇒  $b = d$  and there are  $d^n$  leaves
- Depth-first search for CSPs with single-variable assignments is called *backtracking* search
- Backtracking search is the basic uninformed algorithm for CSPs
- Can solve  $n$ -queens for  $n \approx 25$



## Backtracking search



```
function BACKTRACKING-SEARCH(csp) returns solution/failure
    return RECURSIVE-BACKTRACKING([], csp)
function RECURSIVE-BACKTRACKING(assigned, csp) returns solution/failure
    if assigned is complete then return assigned
    var  $\leftarrow$  SELECT-UNASSIGNED-VARIABLE(VARIABLES[csp], assigned, csp)
    for each value in ORDER-DOMAIN-VALUES(var, assigned, csp) do
        if value is consistent with assigned according to CONSTRAINTS[csp] then
            result  $\leftarrow$  RECURSIVE-BACKTRACKING([var = value | assigned], csp)
            if result  $\neq$  failure then return result
    end
    return failure
```

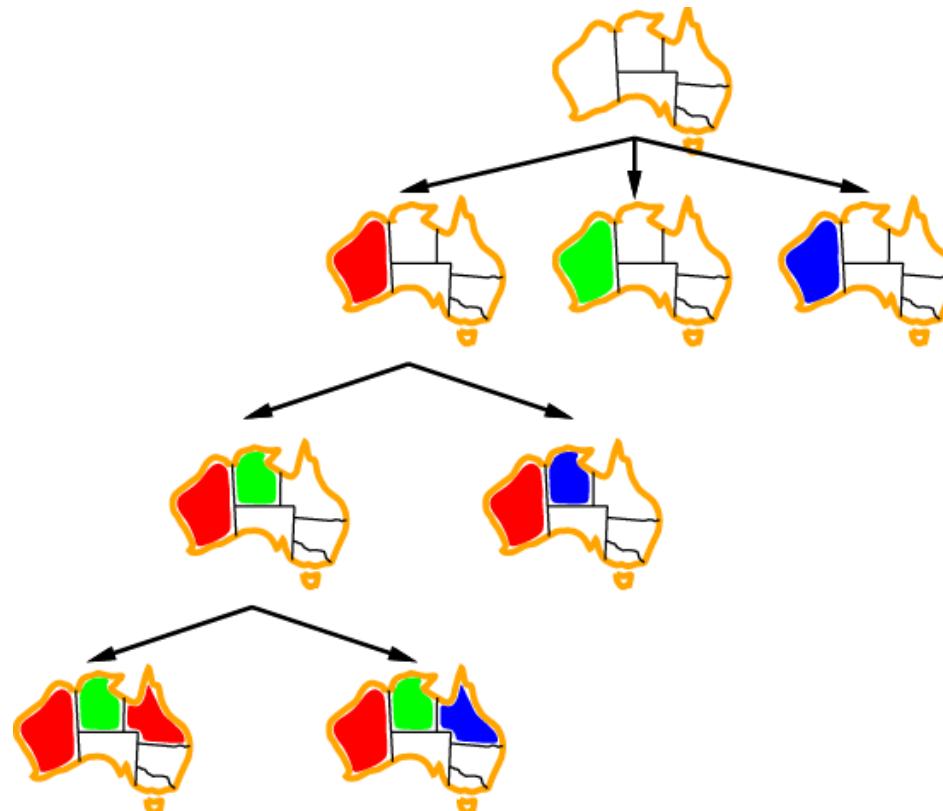


## Backtracking example





## Backtracking example





## Improving backtracking efficiency



*General-purpose* methods can give huge gains in speed:

1. Which variable should be assigned next?
2. In what order should its values be tried?
3. Can we detect inevitable failure early?
4. Can we take advantage of problem structure?



## Most constrained variable



Most constrained variable:

- choose the variable with the fewest legal values

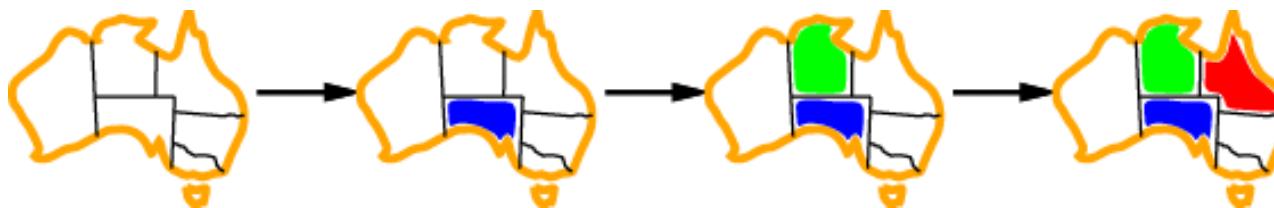




## Most constraining variable



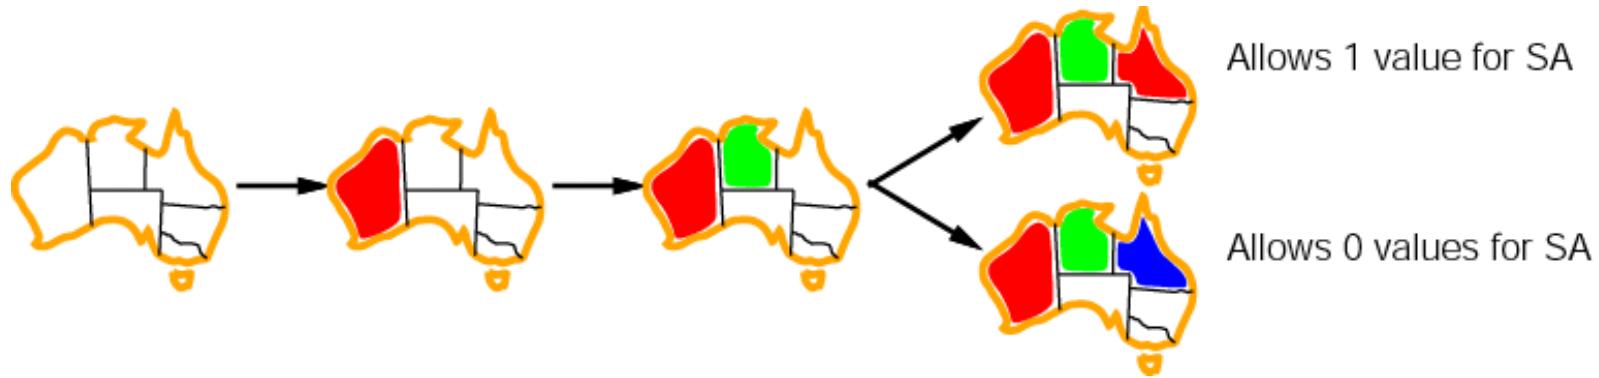
- Tie-breaker among most constrained variables
- Most constraining variable:
  - choose the variable with the most constraints on remaining variables





## Least constraining value

- Given a variable, choose the least constraining value:
  - the one that rules out the fewest values in the remaining variables



Combining these heuristics makes 1000 queens feasible



## Forward checking



**Idea:** Keep track of remaining legal values for unassigned variables

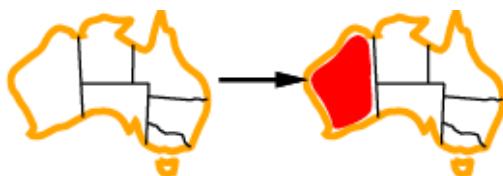
Terminate search when any variable has no legal values



WA	NT	Q	NSW	V	SA	T
█	█	█	█	█	█	█

# Forward checking

**Idea:** Keep track of remaining legal values for unassigned variables  
Terminate search when any variable has no legal values



Region	Total Red	Total Green	Total Blue
WA	1	2	1
NT	1	2	1
Q	1	2	1
NSW	1	2	1
V	1	2	1
SA	1	2	1
T	1	2	1

Region	Red	Green	Blue
WA	1	2	1
NT	0	2	1
Q	0	2	1
NSW	0	2	1
V	0	2	1
SA	0	2	1
T	0	2	1



## Forward checking



**Idea:** Keep track of remaining legal values for unassigned variables

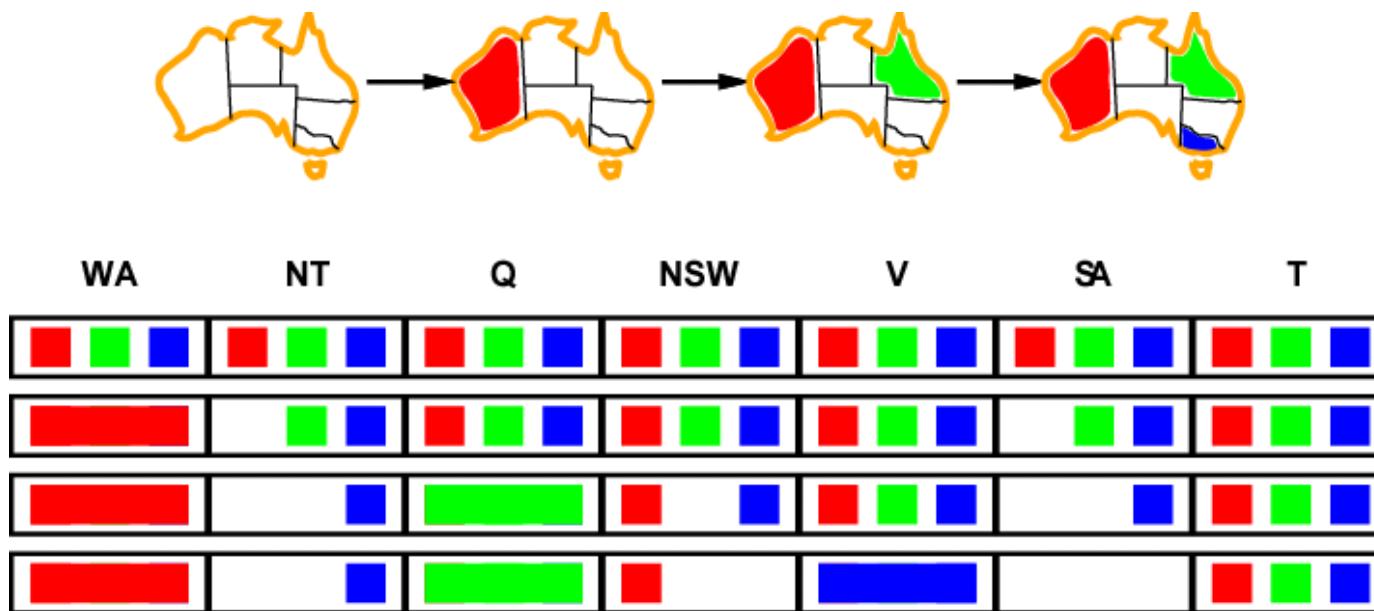
Terminate search when any variable has no legal values



WA	NT	Q	NSW	V	SA	T
█	█	█	█	█	█	█
█		█	█	█	█	█
█			█	█		█

# Forward checking

**Idea:** Keep track of remaining legal values for unassigned variables  
Terminate search when any variable has no legal values



## Constraint propagation

# Constraint propagation

Forward checking propagates information from assigned to unassigned variables, but doesn't provide early detection for all failures:



WA	NT	Q	NSW	V	SA	T
Red	Green	Blue	Red	Green	Blue	Red
Red		Green	Blue	Red	Green	Blue
Red		Blue	Green	Red	Blue	Red

NT and SA cannot both be blue!

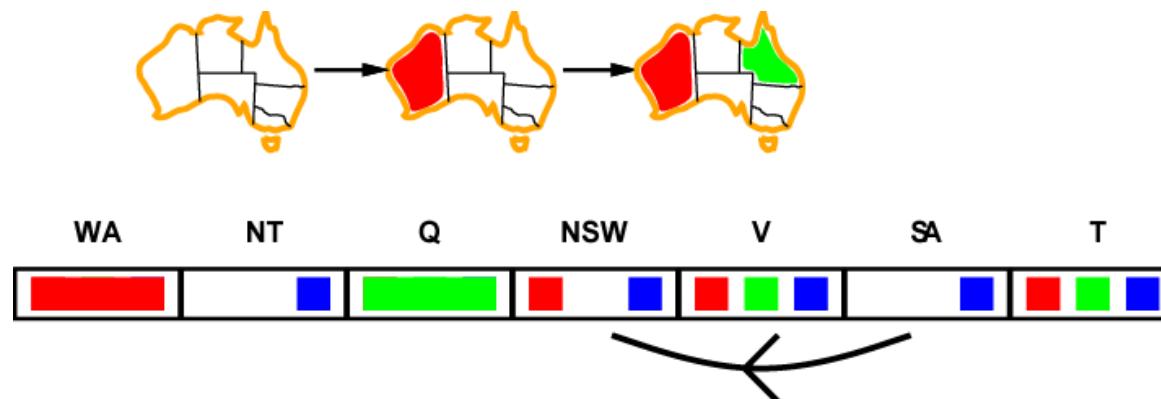
*Constraint propagation* repeatedly enforces constraints locally



## Arc consistency



- Simplest form of propagation makes each arc *consistent*
- $X \rightarrow Y$  is consistent iff
  - for every value  $x$  of  $X$  there is *some* allowed  $y$





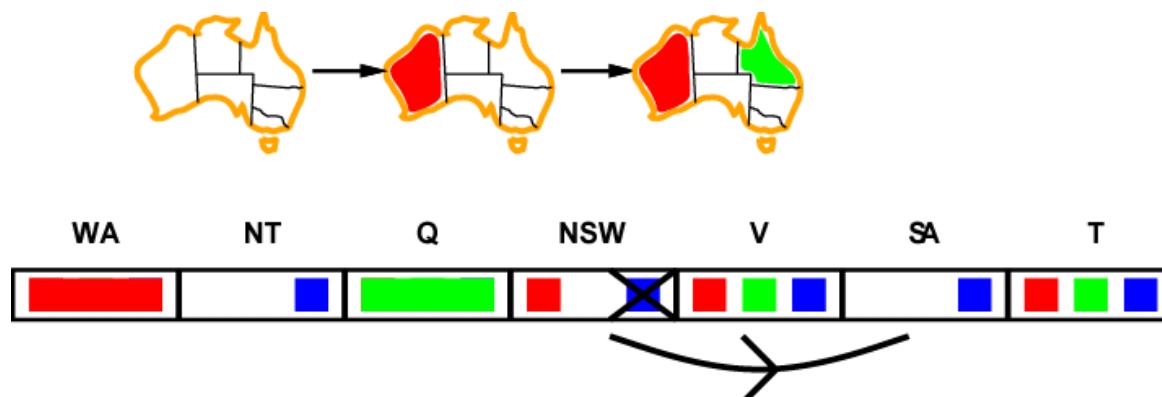
## Arc consistency



Simplest form of propagation makes each arc *consistent*

$X \rightarrow Y$  is consistent iff

- for every value  $x$  of  $X$  there is *some* allowed  $y$





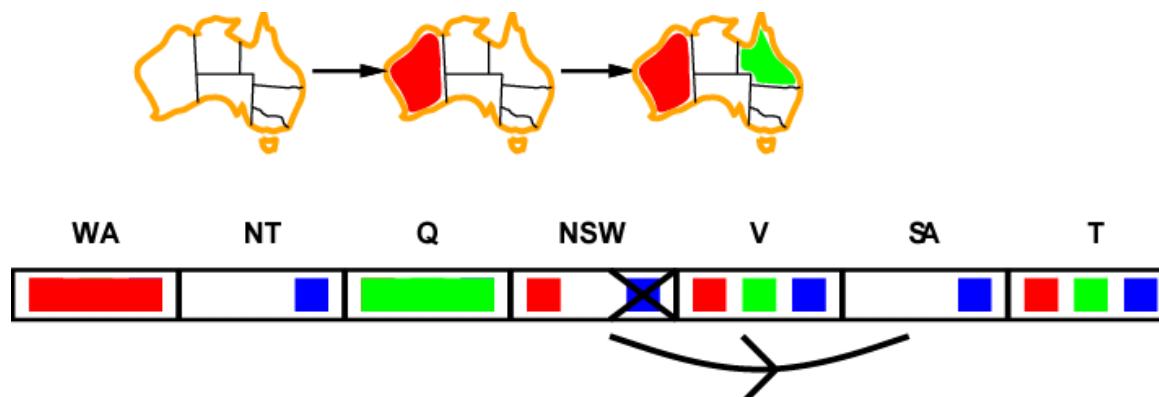
## Arc consistency



Simplest form of propagation makes each arc *consistent*

$X \rightarrow Y$  is consistent iff

- for every value  $x$  of  $X$  there is *some* allowed  $y$





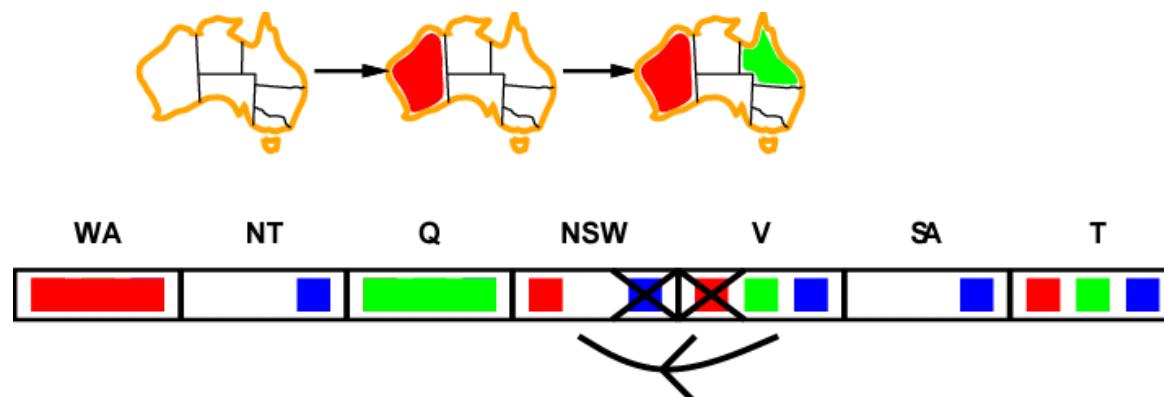
## Arc consistency



Simplest form of propagation makes each arc *consistent*

$X \rightarrow Y$  is consistent iff

- for every value  $x$  of  $X$  there is some allowed  $y$



If  $X$  loses a value, neighbors of  $X$  need to be rechecked



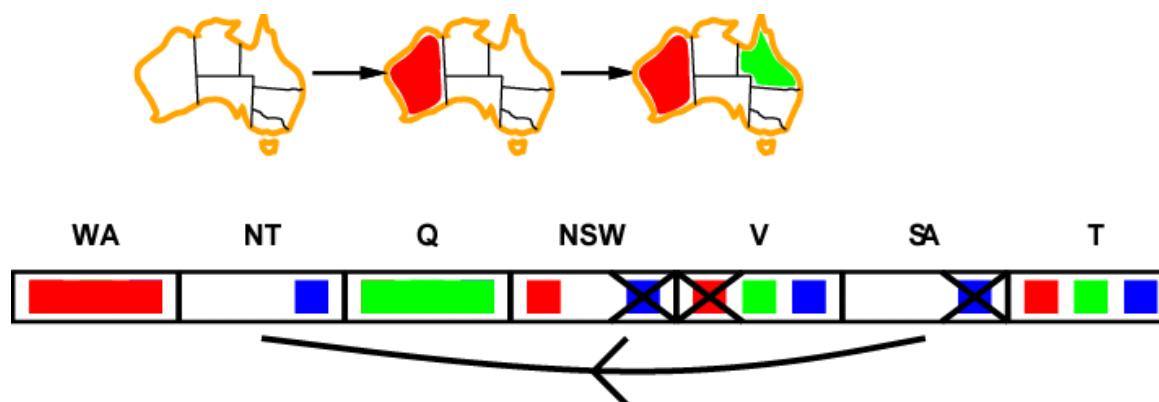
## Arc consistency



Simplest form of propagation makes each arc *consistent*

$X \rightarrow Y$  is consistent iff

- for every value  $x$  of  $X$  there is some allowed  $y$



If  $X$  loses a value, neighbors of  $X$  need to be rechecked

Arc consistency detect failure earlier than forward checking

Can be run as a preprocessor or after each assignment



# Arc consistency algorithm

---

```
function AC-3( csp) returns the CSP, possibly with reduced domains
  inputs: csp, a binary CSP with variables { $X_1, X_2, \dots, X_n$ }
  local variables: queue, a queue of arcs, initially all the arcs in csp

  while queue is not empty do
    ( $X_i, X_j$ )  $\leftarrow$  REMOVE-FIRST(queue)
    if REMOVE-INCONSISTENT-VALUES( $X_i, X_j$ ) then
      for each  $X_k$  in NEIGHBORS[ $X_i$ ] do
        add ( $X_k, X_i$ ) to queue

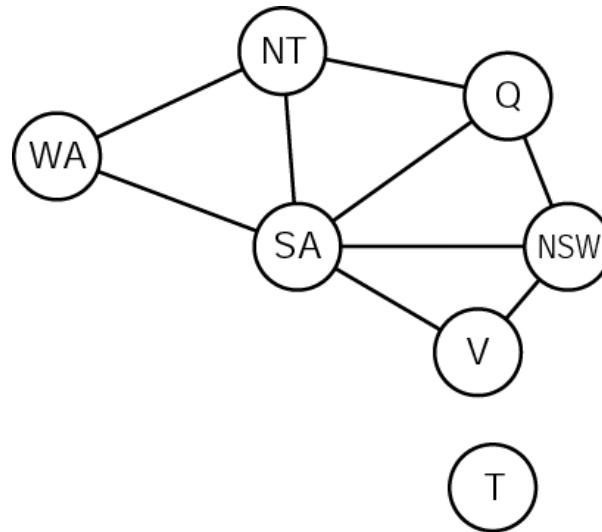

---


function REMOVE-INCONSISTENT-VALUES(  $X_i, X_j$ ) returns
true iff we remove a value
  removed  $\leftarrow$  false
  for each  $x$  in DOMAIN[ $X_i$ ] do
    if no value  $y$  in DOMAIN[ $X_j$ ] allows  $(x,y)$  to satisfy constraint between  $X_i$  and  $X_j$ 
      then delete  $x$  from DOMAIN[ $X_i$ ]; removed  $\leftarrow$  true
  return removed
```

---

- $O(n^2 d^3)$ , can be reduced to  $O(n^2 d^2)$  but cannot detect all failures in poly time!

# Problem structure



- Tasmania and mainland are *independent subproblems*
- Identifiable as *connected components* of constraint graph



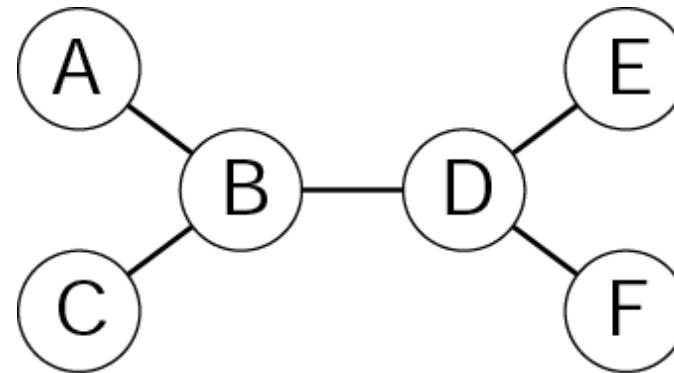
## Problem structure (contd.)



- Suppose each subproblem has  $c$  variables out of  $n$  total
- Worst-case solution cost is  $n/c * d^c$ , *linear* in  $n$
- E.g.,  $n = 80$ ,  $d = 2$ ,  $c = 20$ 
  - $2^{80} = 4$  billion years at 10 million nodes/sec
  - $4 * 2^{20} = 0.4$  seconds at 10 million nodes/sec



## Tree-structured CSPs

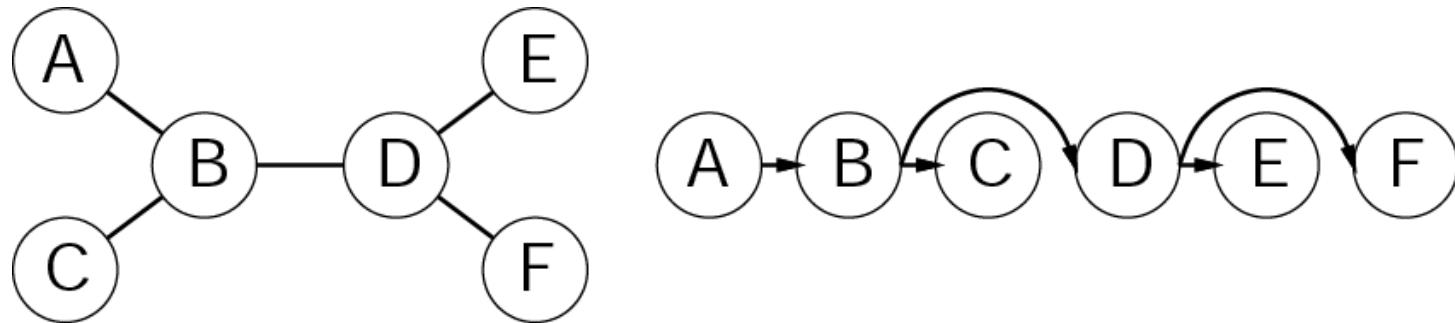


- **Theorem:** if the constraint graph has no loops, the CSP can be solved in  $O(n d^2)$  time
- Compare to general CSPs, where worst-case time is  $O(d^n)$
- This property also applies to logical and probabilistic reasoning: an important example of the relation between syntactic restrictions and the complexity of reasoning.



## Algorithm for tree-structured CSPs

1. Choose a variable as root, order variables from root to leaves such that every node's parent precedes it in the ordering

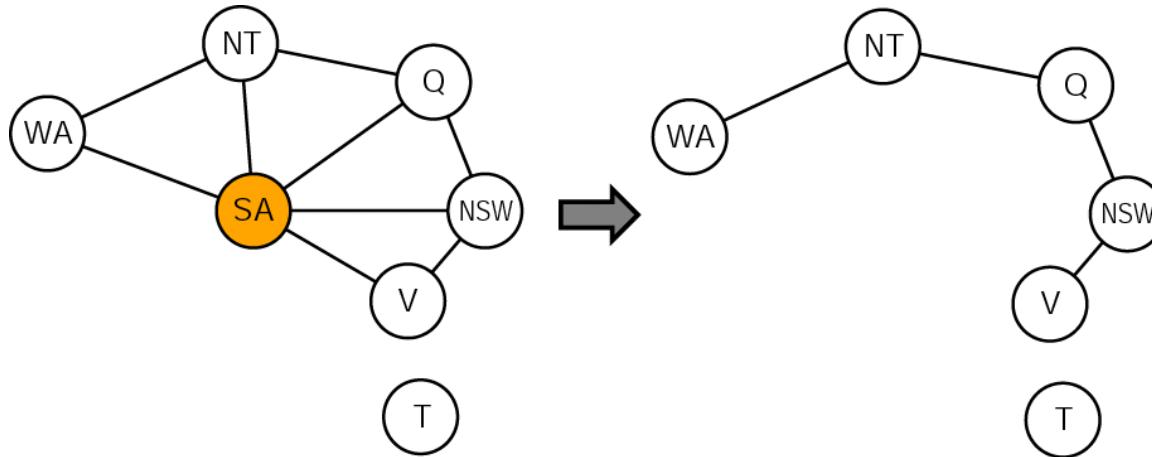


2. For  $j$  from  $n$  down to 2, apply **RemoveInconsistent**( $\text{Parent}(X_j), X_j$ )
3. For  $j$  from 1 to  $n$ , assign  $X_j$  consistently with  $\text{Parent}(X_j)$



## Nearly tree-structured CSPs

- **Conditioning:** instantiate a variable, prune its neighbors' domains



- **Cutset conditioning:** instantiate (in all ways) a set of variables such that the remaining constraint graph is a tree
- Cutset size  $\Rightarrow$  runtime  $O(d^c \bullet (n-c)d^2)$ , very fast for small c



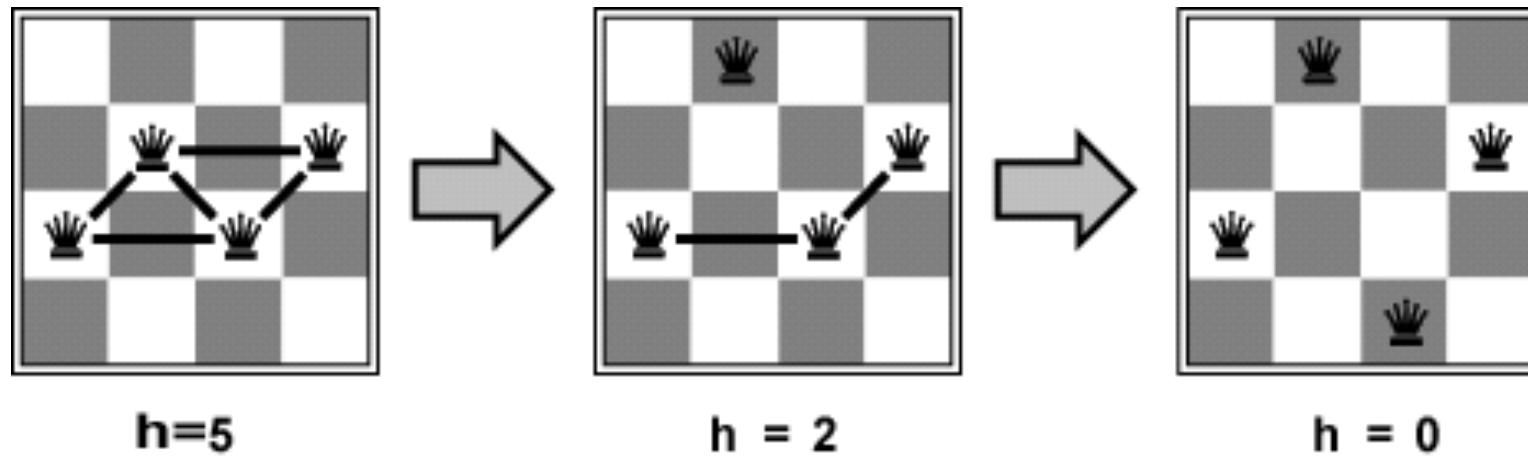
## Iterative algorithms for CSPs

- Hill-climbing, simulated annealing typically work with “complete” states, i.e., all variables assigned
- To apply to CSPs:
  - allow states with unsatisfied constraints
  - operators *reassign* variable values
- Variable selection: randomly select any conflicted variable
- Value selection by *min-conflicts* heuristic:
  - choose value that violates the fewest constraints
  - i.e., hillclimb with  $h(n)$  = total number of violated constraints



## Example: 4-Queens

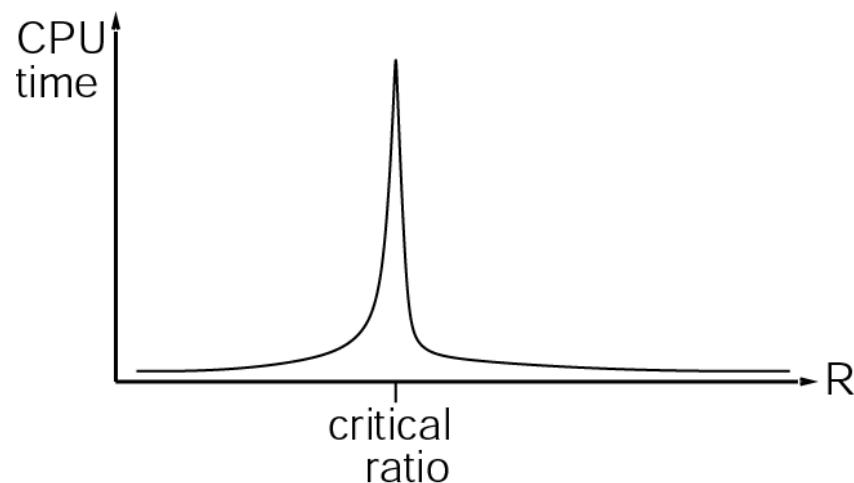
- **States:** 4 queens in 4 columns ( $4^4 = 256$  states)
- **Operators:** move queen in column
- **Goal test:** no attacks
- **Evaluation:**  $h(n) = \text{number of attacks}$





## Performance of min-conflicts

- Given random initial state, can solve  $n$ -queens in almost constant time for arbitrary  $n$  with high probability (e.g.,  $n = 10,000,000$ )
- The same appears to be true for any randomly-generated CSP *except* in a narrow range of the ratio:  $R = \frac{\text{number of constraints}}{\text{number of variables}}$





# Summary

- CSPs are a special kind of problem:
  - states defined by values of a fixed set of variables
  - goal test defined by *constraints* on variable values
- Backtracking = depth-first search with one variable assigned per node
- Variable ordering and value selection heuristics help significantly
- Forward checking prevents assignments that guarantee later failure
- Constraint propagation (e.g., arc consistency) does additional work to constrain values and detect inconsistencies
- The CSP representation allows analysis of problem structure
- Tree-structured CSPs can be solved in linear time
- Iterative min-conflicts is usually effective in practice

# Practical AI (a.k.a. Business Intelligence)

## ***Part 1- Intelligent Search***

Search

(Intelligent) Search

Constraint Satisfaction

**Adversarial Search (Games)**



**Universität  
Zürich<sup>UZH</sup>**



Dynamic and Distributed  
Information Systems





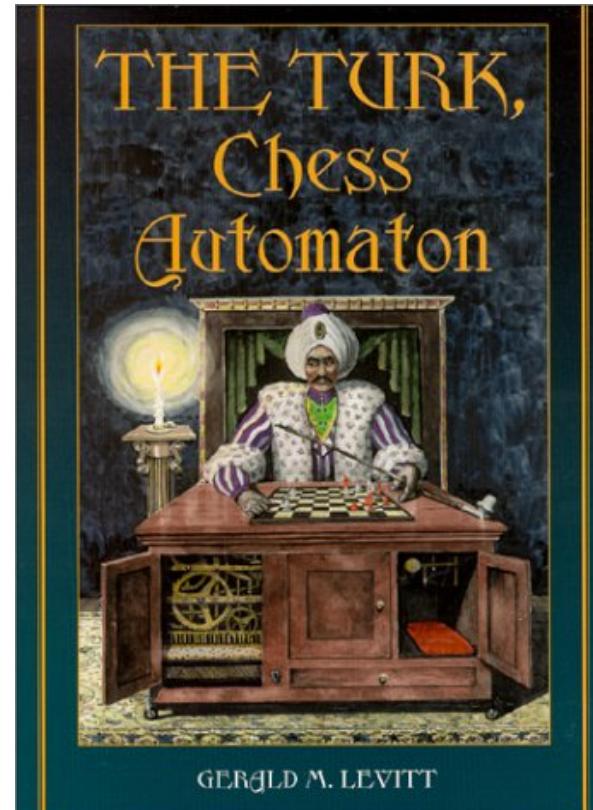
## Overview – Part 1



- Problem Definition?
- Search
  - Breadth first, Depth first, Iterative deepening
  - Best First
- Informed (or Intelligent) Search
  - A\*, Hill climbing, Simulated annealing
- Constraint Satisfaction
  - Forward checking, Constraint propagation
- **Adversarial Search (Games)**

# Outline

- Perfect play
- Resource limits
- $\alpha-\beta$  pruning
- Games of chance
- Games of imperfect information





## Games vs. search problems



- „Unpredictable“ opponent → solution is a strategy specifying a move for every possible opponent reply
- Time limits → unlikely to find goal, must approximate
- Plan of attack:
  - Computer considers possible lines of play (Babbage, 1846)
  - Algorithm for perfect play (Zermelo, 1912; Von Neumann, 1944)
  - Finite horizon, approximate evaluation (Zuse, 1945; Wiener, 1948; Shannon, 1950)
  - First chess program (Turing, 1951)
  - Machine learning to improve evaluation accuracy (Samuel, 1952-57)
  - Pruning to allow deeper search (McCarthy, 1956)



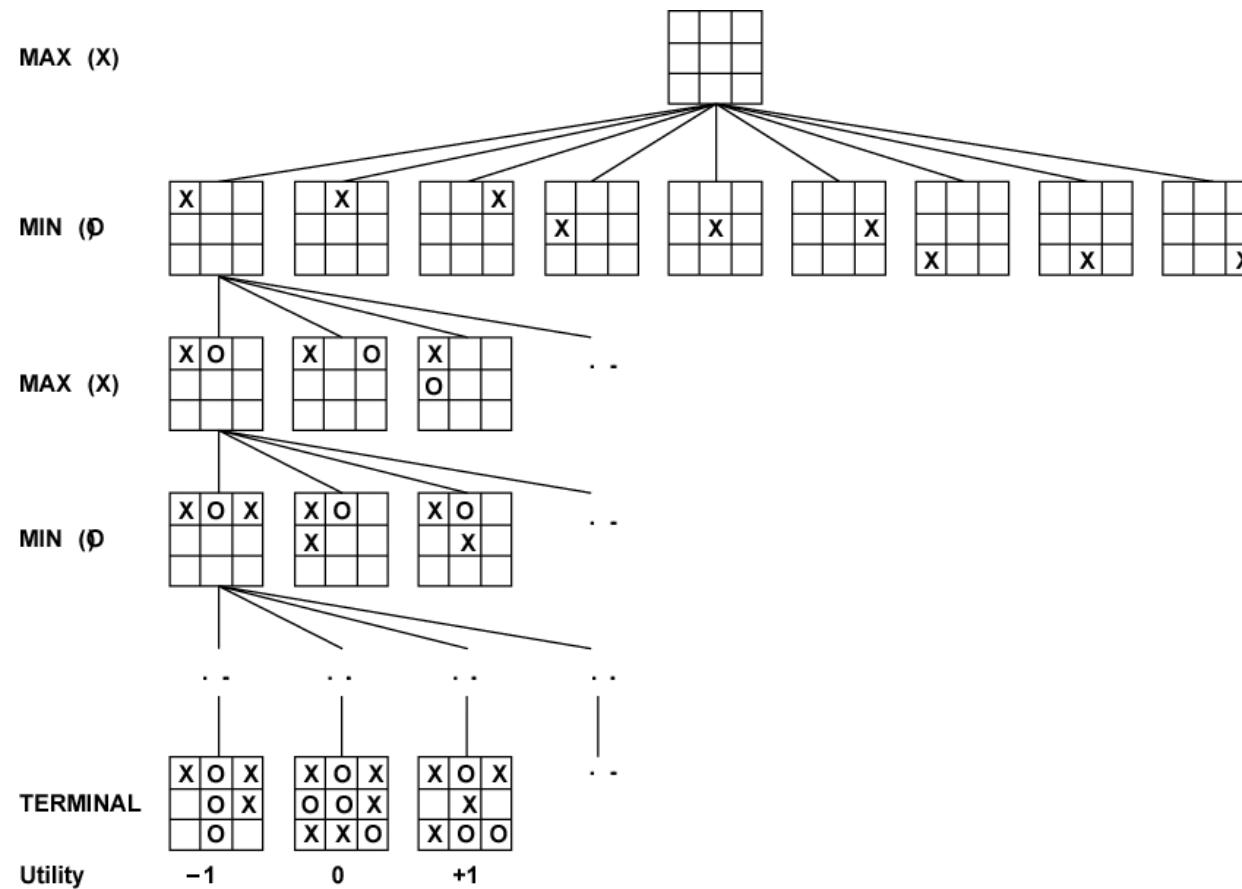
## Types of games



	<i>Deterministic</i>	<i>Chance</i>
<i>Perfect information</i>	Chess, checkers, go, othello	Backgammon, monopoly
<i>Imperfect information</i>		Bridge, poker, scrabble, nuclear war



# Game tree (2-player, deterministic, turns)

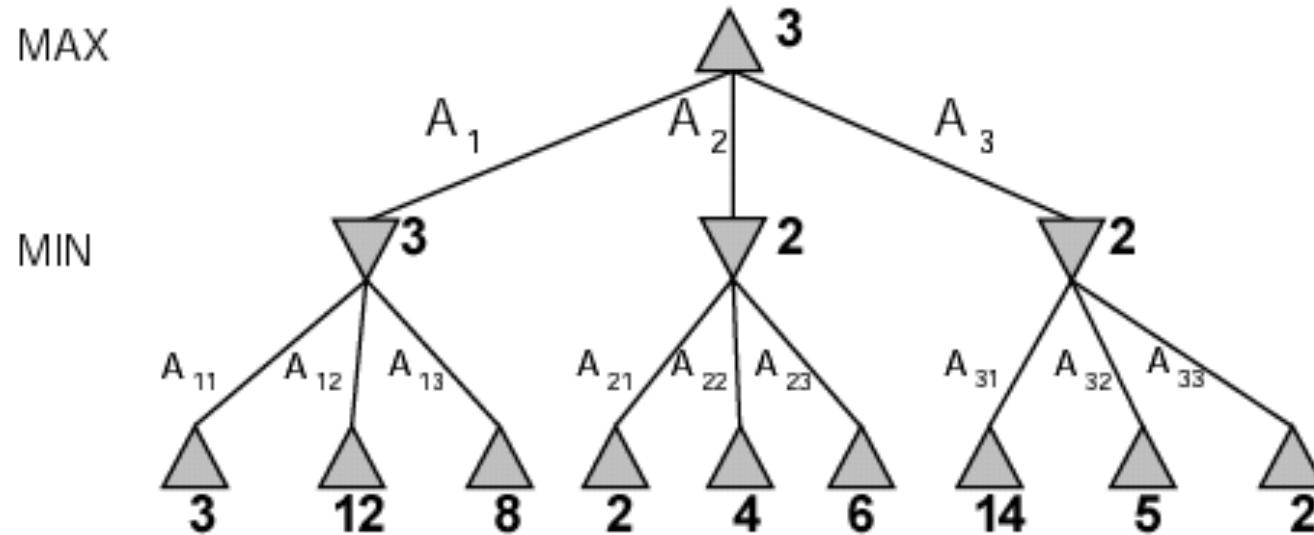




## Minimax



- Perfect play for deterministic, perfect-information games
- **Idea:** choose move to position with highest *minimax value*  
= best achievable payoff against best play
- E.g., 2-ply game:





## Minimax algorithm



**function** MINIMAX-DECISION(*state, game*) **returns** *an action*

*action, state*  $\leftarrow$  the *a, s* **in** SUCCESSORS(*state*)

such that MINIMAX-VALUE(*s, game*) is maximized

**return** *action*

---

**function** MINIMAX-VALUE(*state, game*) **returns** *a utility value*

**if** TERMINAL-TEST(*state*) **then**

**return** UTILITY(*state*)

**else if** MAX is to move in *state* **then**

**return** the highest MINIMAX-VALUE of SUCCESSORS(*state*)

**else**

**return** the lowest MINIMAX-VALUE of SUCCESSORS(*state*)



## Properties of minimax



- **Complete?**
  - Yes, if tree is finite (chess has specific rules for this)
- **Optimal?**
  - Yes, against an optimal opponent. Otherwise??
- **Time complexity?**
  - $O(b^m)$
- **Space complexity?**
  - $O(bm)$  (depth-first exploration)
- For chess,  $b \approx 35$ ,  $m \approx 100$  for „reasonable“ games
  - exact solution completely infeasible



## Resource limits



- Suppose we have 100 seconds, explore  $10^4$  nodes/second  
→  $10^6$  nodes per move
- Standard approach:
  - *cutoff test*  
e.g., depth limit (perhaps add *quiescence search*)
  - *evaluation function*  
= estimated desirability of position



## Evaluation functions



- For chess, typically *linear* weighted sum of *features*  
 $\text{Eval}(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$
- e.g.,  $w_1 = 9$  with  
 $f_1(s) = (\text{number of white queens}) - (\text{number of black queens})$ , etc.



(a) White to move  
Fairly even



(b) Black to move  
White slightly better



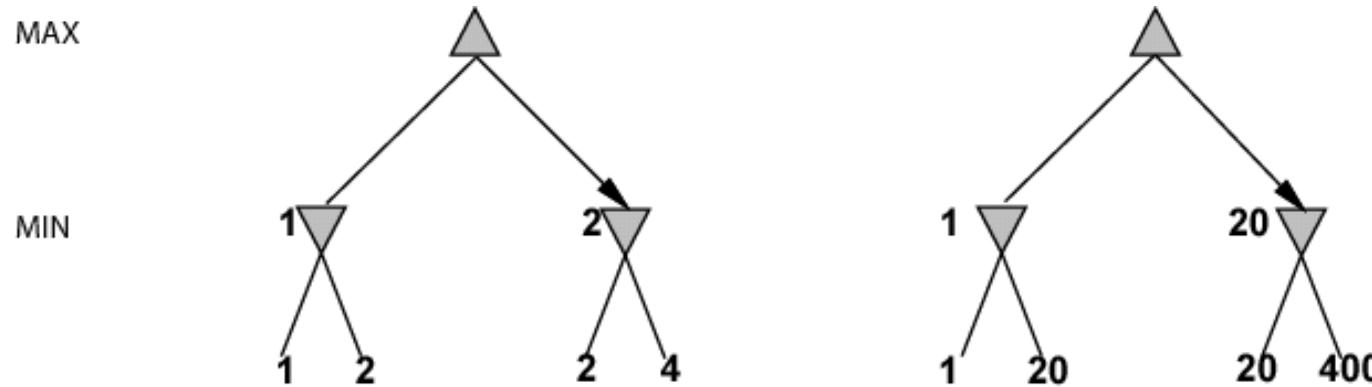
(c) White to move  
Black winning



(d) Black to move  
White about to lose



## Digression: Exact values don't matter



- Behaviour is preserved under any *monotonic* transformation of **Eval**
- Only the order matters:  
payoff in deterministic games acts as an *ordinal utility* function



## Cutting off search

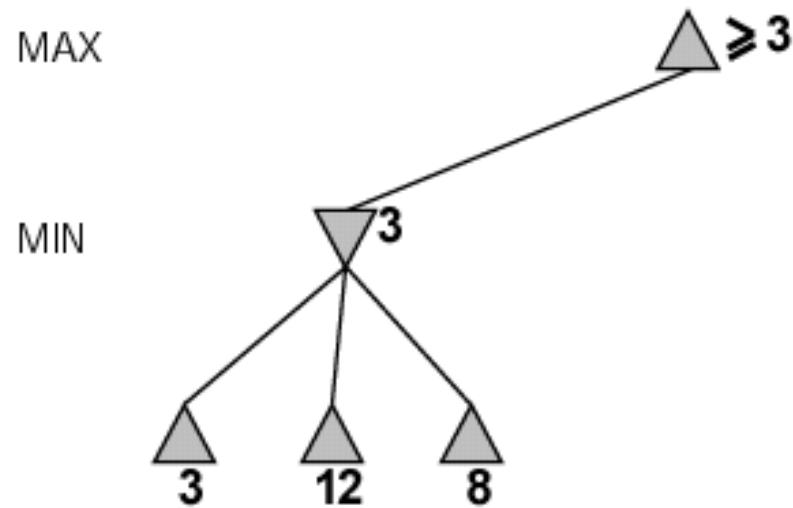


- **MinimaxCutoff** is identical to **MinimaxValue** except:
  1. **Terminal?** is replaced by **Cutoff?**
  2. **Utility** is replaced by **Eval**
- Does it work in practice?
  - $b^m = 10^6$ ,  $b=35 \rightarrow m=4$
  - 4-ply lookahead is a hopeless chess player!
  - 4-ply  $\approx$  human novice
  - 8-ply  $\approx$  typical PC, human master
  - 12-ply  $\approx$  Deep Blue, Kasparov

Γ

## α-β pruning example

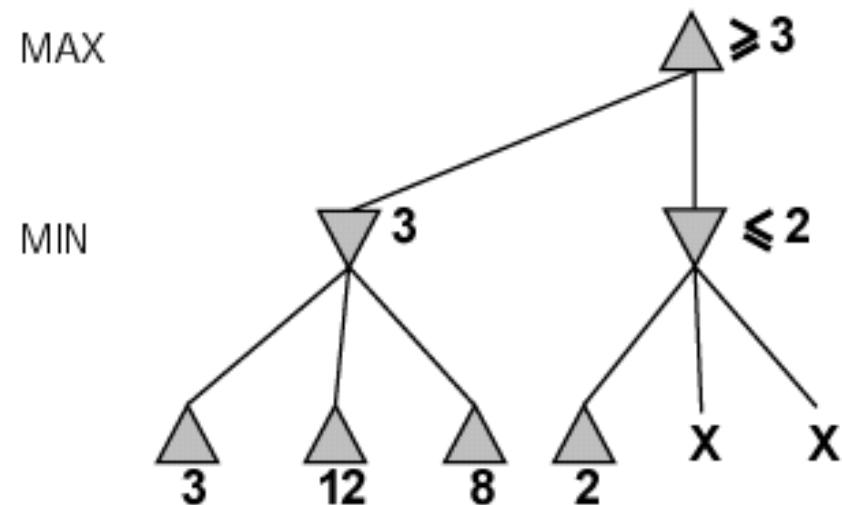
└



[

## $\alpha$ - $\beta$ pruning example

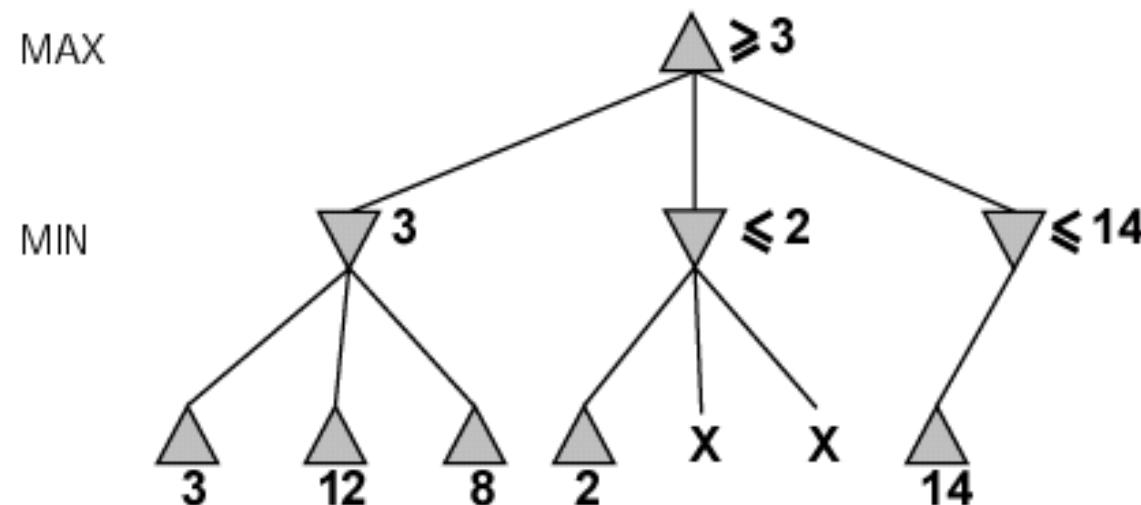
]



[

## $\alpha$ - $\beta$ pruning example

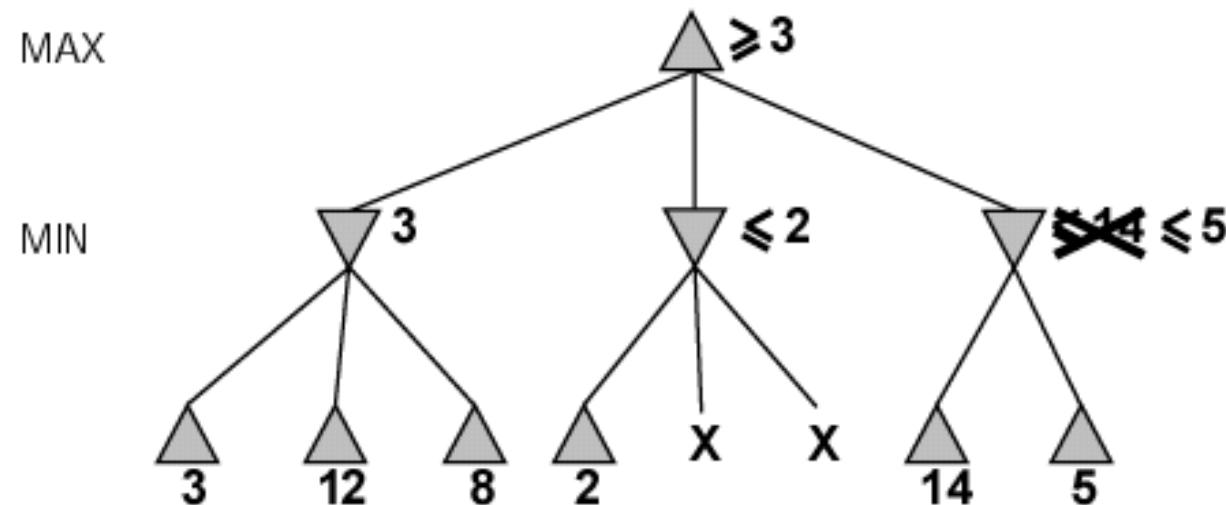
]



[

## $\alpha$ - $\beta$ pruning example

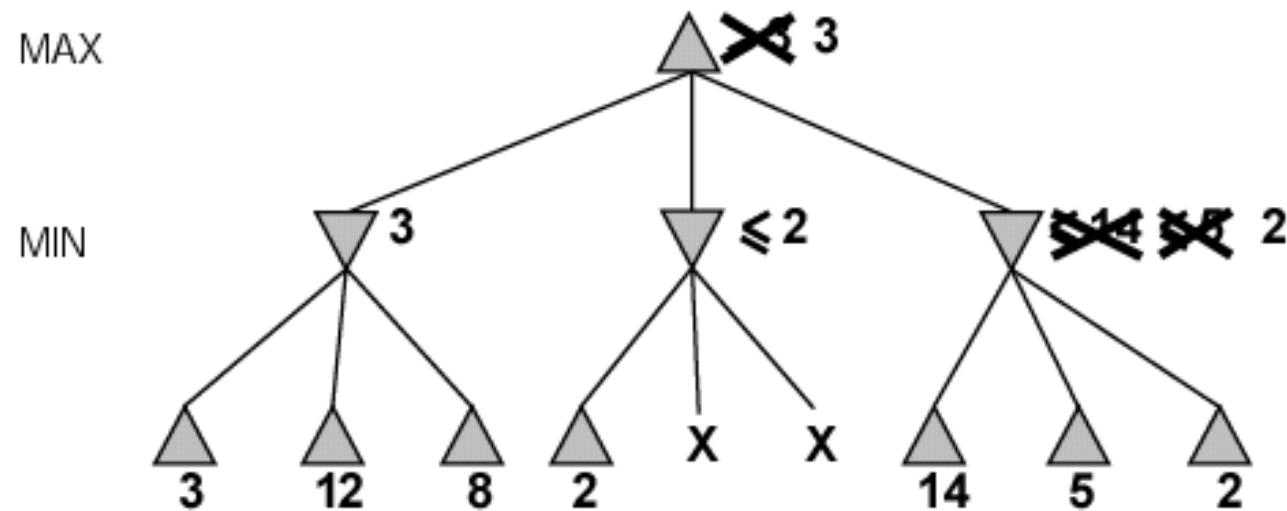
]



[

## $\alpha$ - $\beta$ pruning example

]





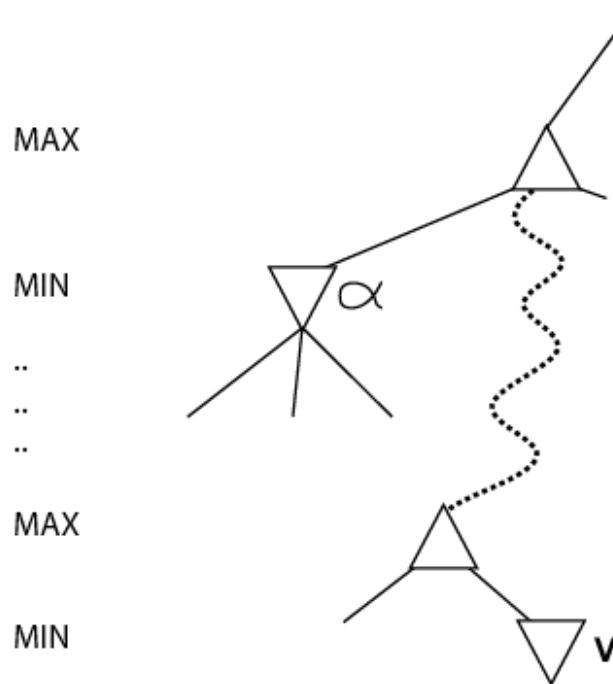
## Properties of $\alpha$ - $\beta$



- Pruning *does not* affect final result
- Good move ordering improves effectiveness of pruning
- With „perfect ordering,” time complexity =  $O(b^{m/2})$ 
  - *doubles* depth of search
  - can easily reach depth 8 and play good chess
- A simple example of the value of reasoning about which computations are relevant (a form of *metareasoning*)



## Why is it called $\alpha$ - $\beta$ ?



- $\alpha$  is the best value **Max** found so far off the current path
- If  $V$  is worse than  $\alpha$  **Max** will avoid it  $\Rightarrow$  prune that branch
- Define  $\beta$  similarly for **Min**



## The $\alpha$ - $\beta$ algorithm

**function** ALPHA-BETA-SEARCH(*state, game*) **returns** an action  
    *action, state*  $\leftarrow$  the *a, s* **in** SUCCESSORS[*game*](*state*)  
        such that MIN-VALUE(*s, game, -∞, +∞*) is maximized  
**return** *action*

---

**function** MAX-VALUE(*state, game, α, β*) **returns** the minimax value of *state*  
    **if** CUTOFF-TEST(*state*) **then return** EVAL(*state*)  
    **for each** *s* **in** SUCCESSORS(*state*) **do**  
         $\alpha \leftarrow \max(\alpha, \text{MIN-VALUE}(s, \text{game}, \alpha, \beta))$   
        **if**  $\alpha \geq \beta$  **then return**  $\beta$   
**return**  $\alpha$

---

**function** MIN-VALUE(*state, game, α, β*) **returns** the minimax value of *state*  
    **if** CUTOFF-TEST(*state*) **then return** EVAL(*state*)  
    **for each** *s* **in** SUCCESSORS(*state*) **do**  
         $\beta \leftarrow \min(\beta, \text{MAX-VALUE}(s, \text{game}, \alpha, \beta))$   
        **if**  $\beta \leq \alpha$  **then return**  $\alpha$   
**return**  $\beta$



## Deterministic games in practice



**Checkers**: Chinook ended 40-year-reign of human world champion Marion Tinsley in 1994. Used an endgame database defining perfect play for all positions involving 8 or fewer pieces on the board, a total of 443,748,401,247 positions.

**Chess**: Deep Blue defeated human world champion Gary Kasparov in a six-game match in 1997. Deep Blue searches 200 million positions per second, uses very sophisticated evaluation, and undisclosed methods for extending some lines of search up to 40 ply.

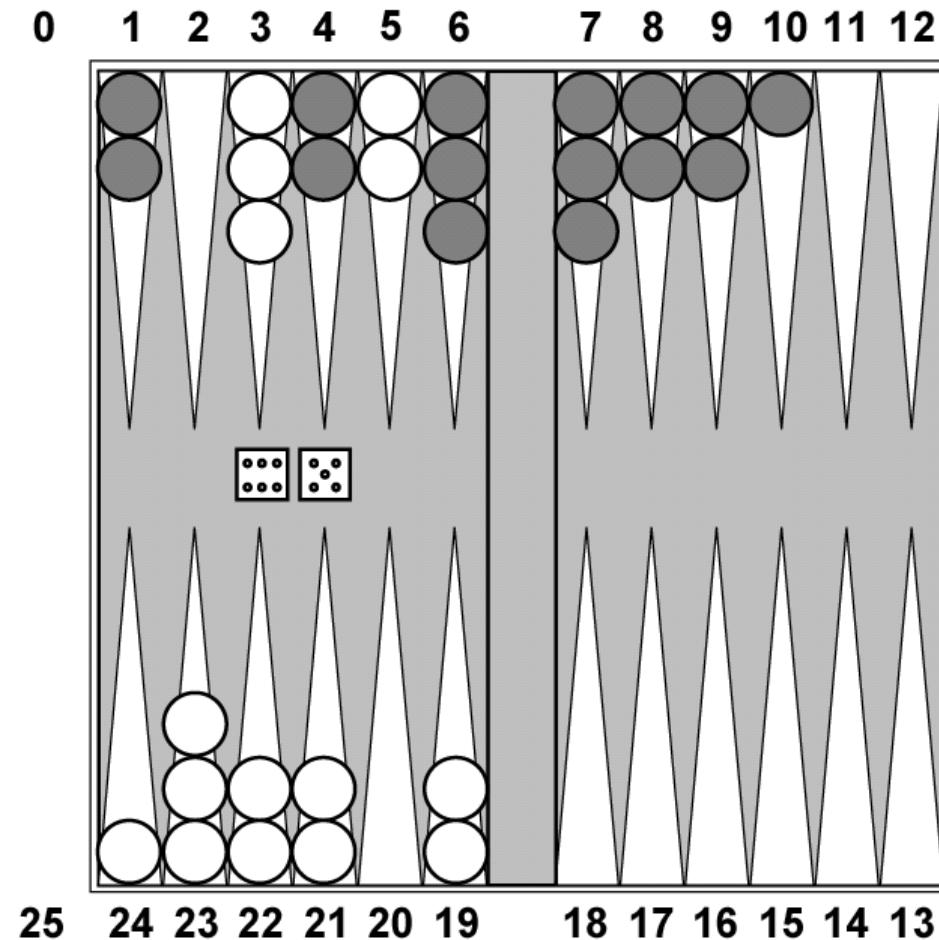
**Othello**: human champions refuse to compete against computers, who are too good.

**Go**: human champions refuse to compete against computers, who are too bad. In go,  $b > 300$ , so most programs use pattern knowledge bases to suggest plausible moves.

What happened here?  
Monte Carlo Tree search and ML



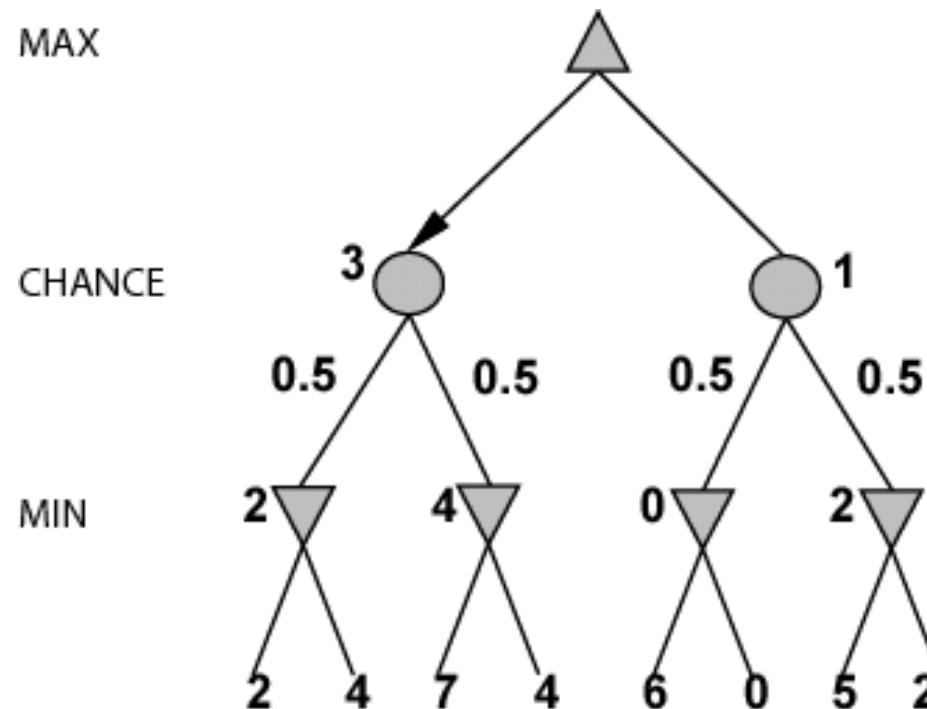
## Nondeterministic games: backgammon



## └

# Nondeterministic games in general

- In nondeterministic games, chance introduced by dice, card-shuffling
- Simplified example with coin-flipping:





# Algorithm for nondeterministic games

- **Expectiminimax** gives perfect play
- Just like **Minimax**, except we must also handle chance nodes:

...

If *state* is a **Max** node then

return the **highest Expectimax-Value of Successors** (*state*)

If *state* is a **Min** node then

return the **lowest Expectimax-Value of Successors** (*state*)

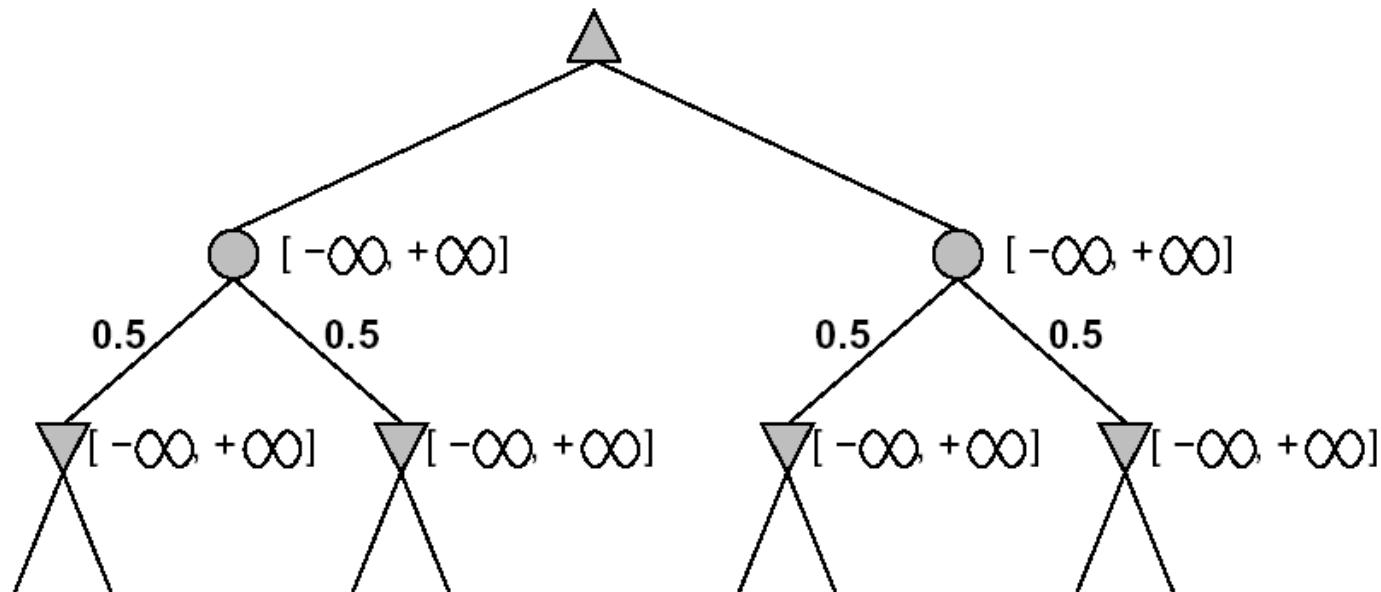
If *state* is a chance node then

return the **average Expectimax-Value of Successors** (*state*)

...

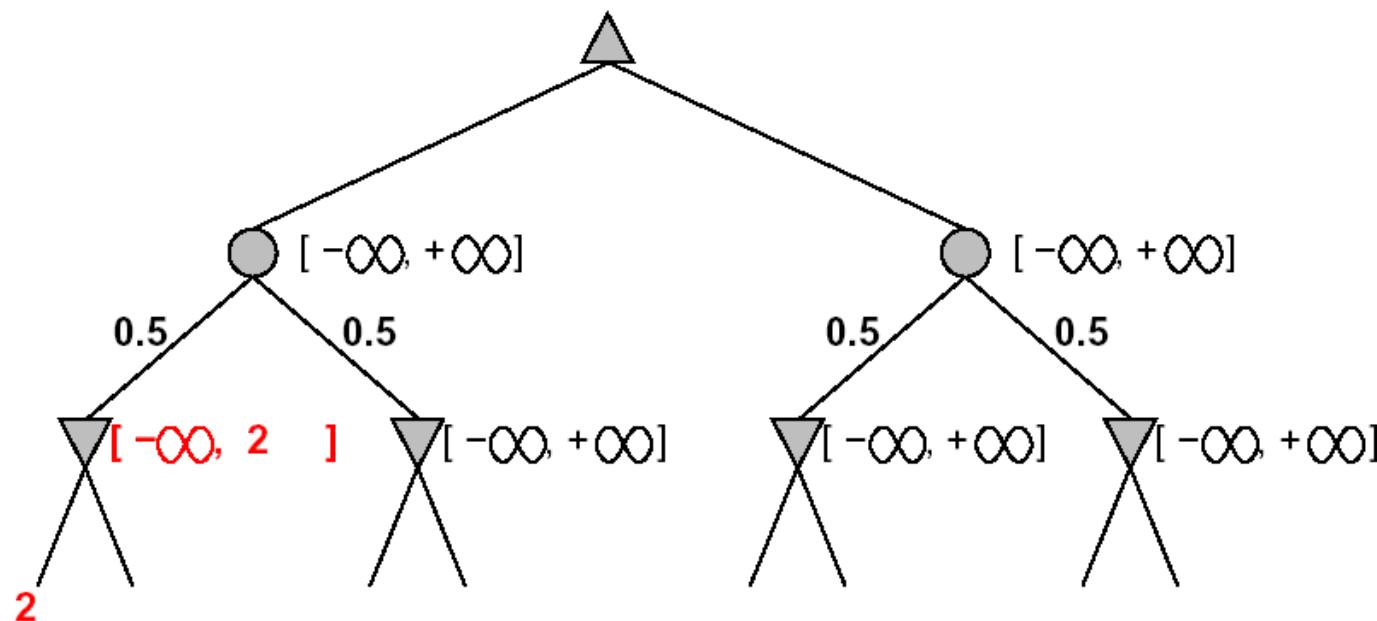
## Pruning in nondeterministic game trees

- A version of  $\alpha$ - $\beta$  pruning is possible:



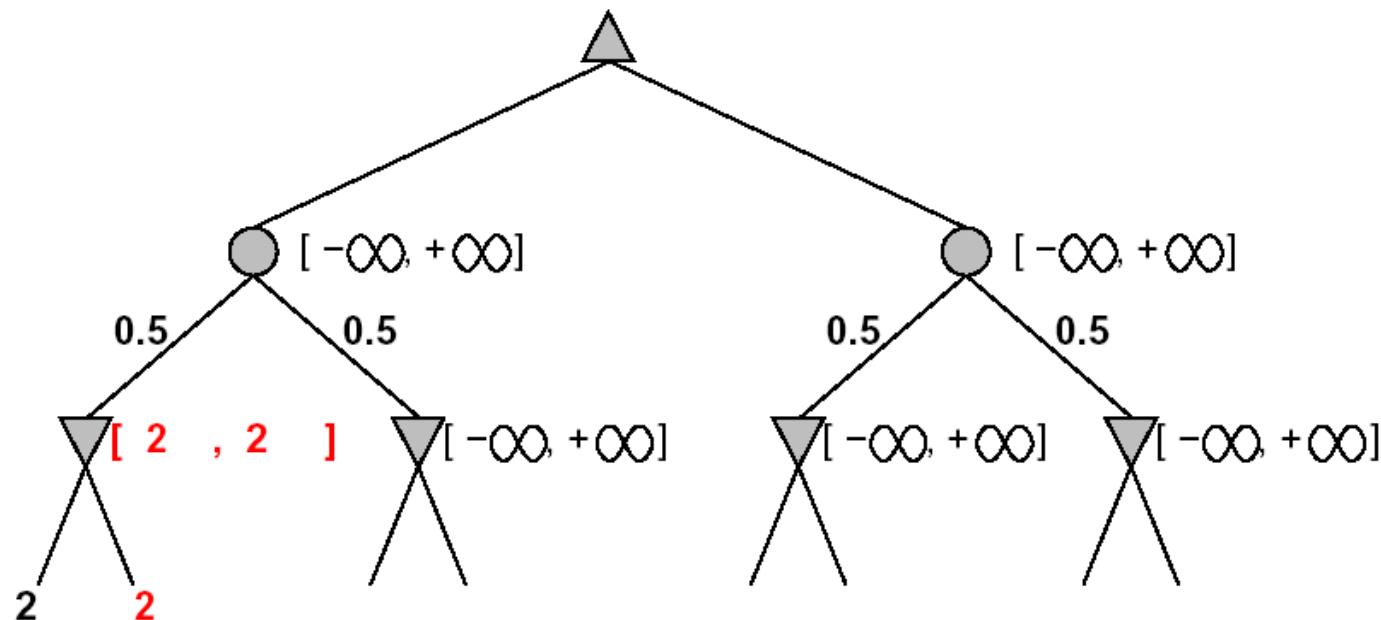
# Pruning in nondeterministic game trees

- A version of  $\alpha$ - $\beta$  pruning is possible:



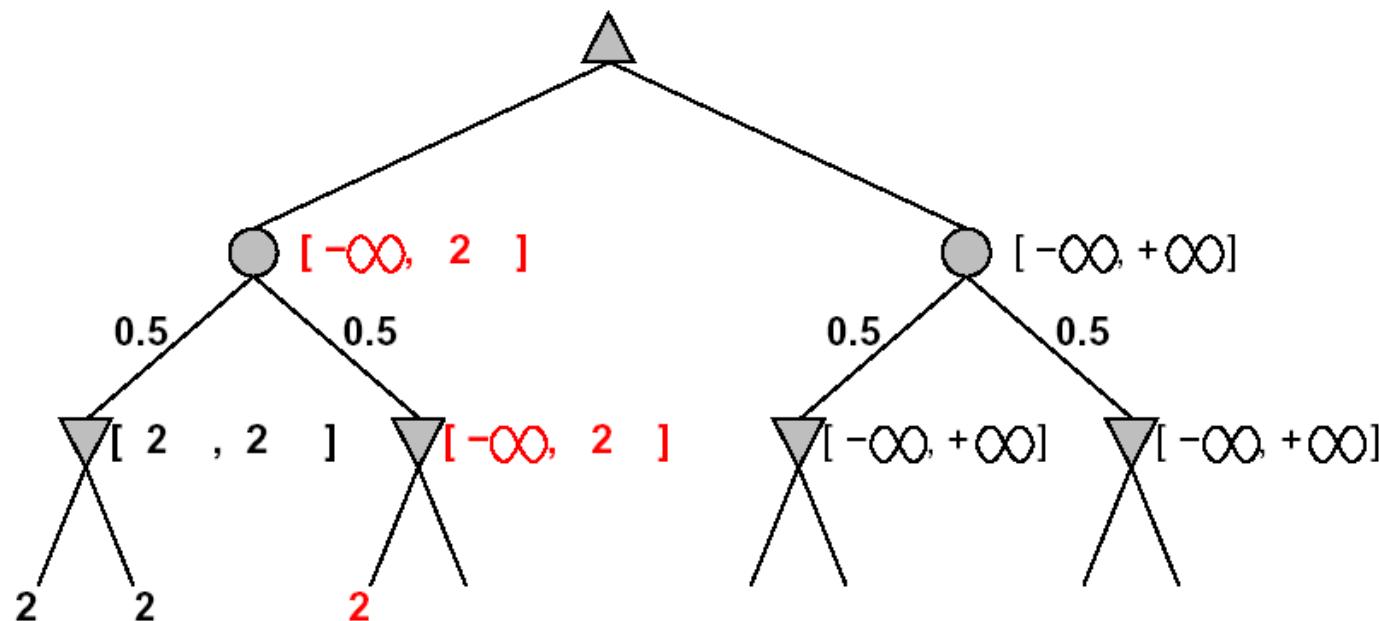
## Pruning in nondeterministic game trees

- A version of  $\alpha$ - $\beta$  pruning is possible:



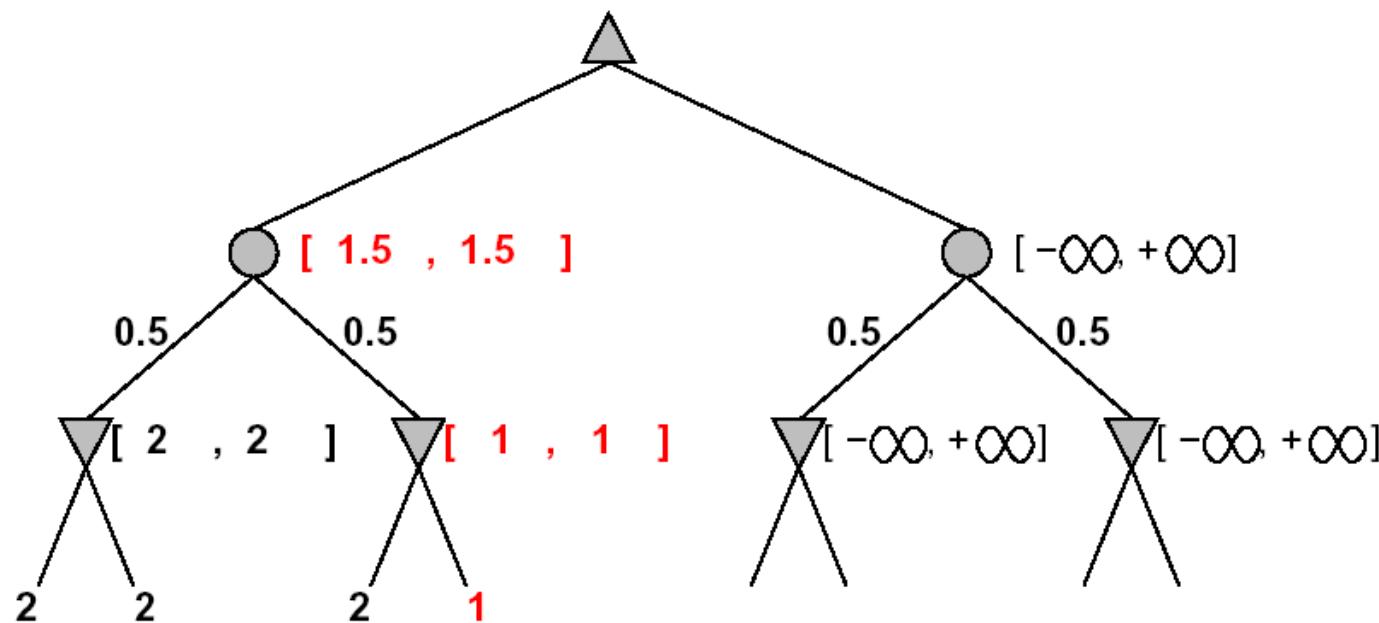
## Pruning in nondeterministic game trees

- A version of  $\alpha$ - $\beta$  pruning is possible:



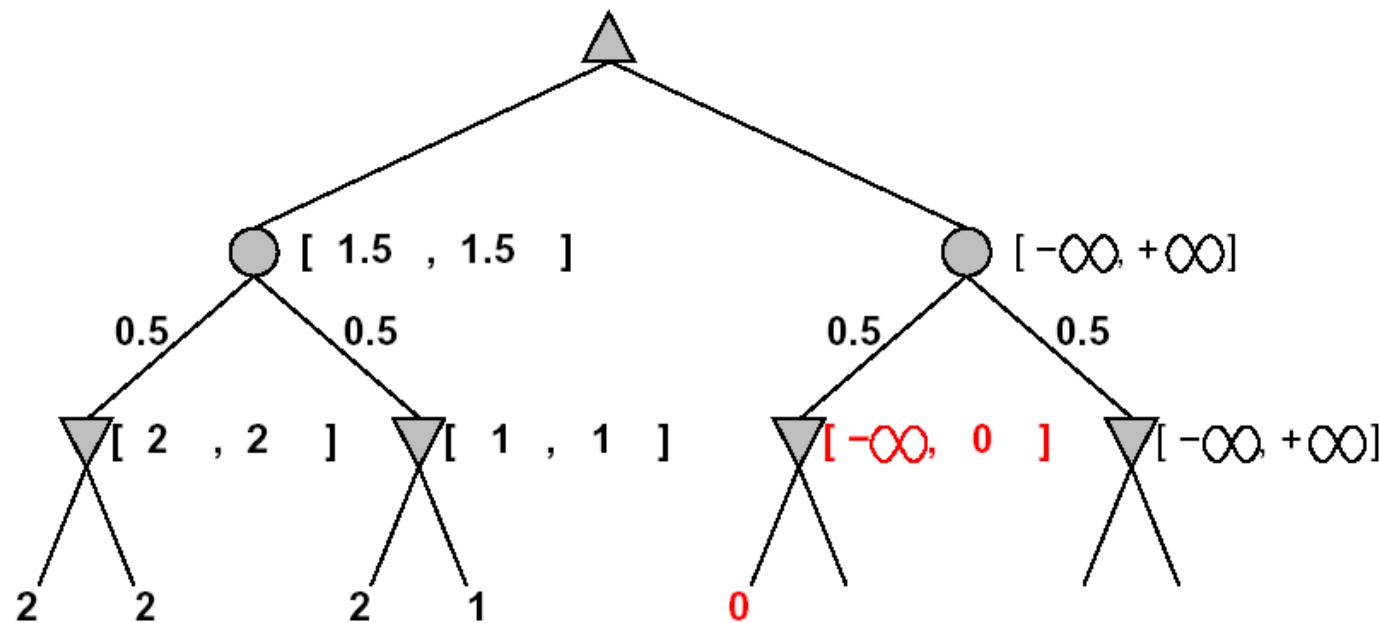
## Pruning in nondeterministic game trees

- A version of  $\alpha$ - $\beta$  pruning is possible:



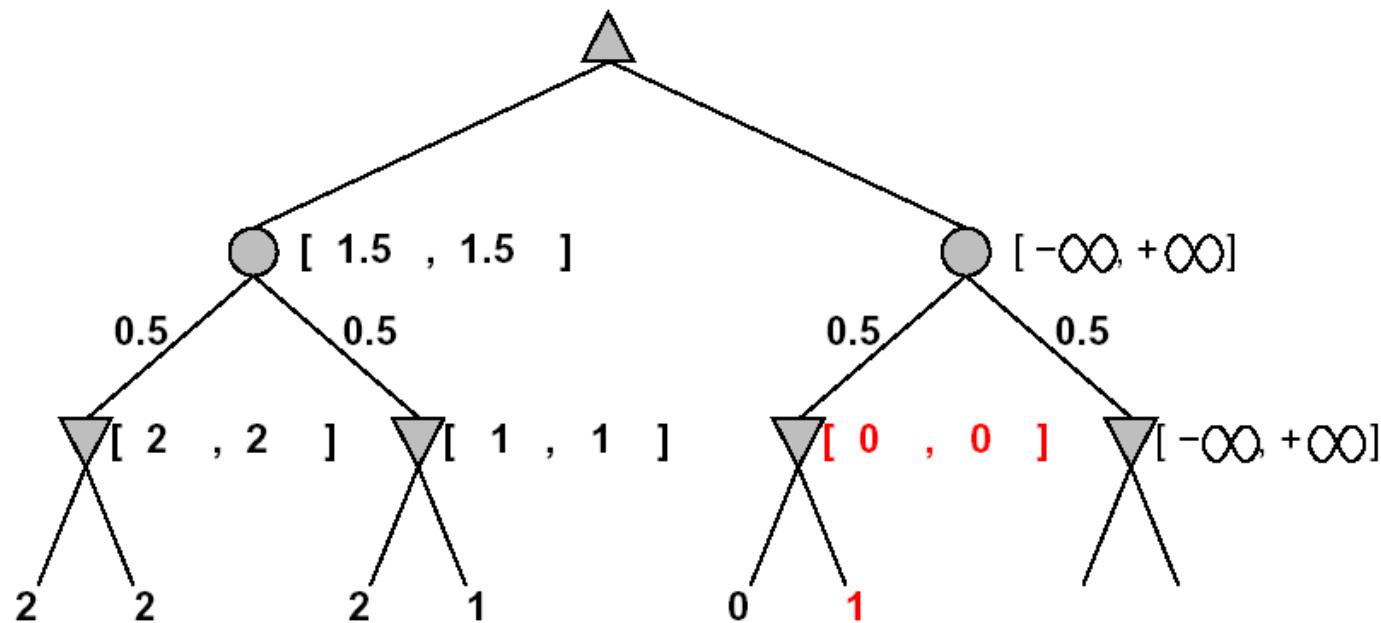
## Pruning in nondeterministic game trees

- A version of  $\alpha$ - $\beta$  pruning is possible:



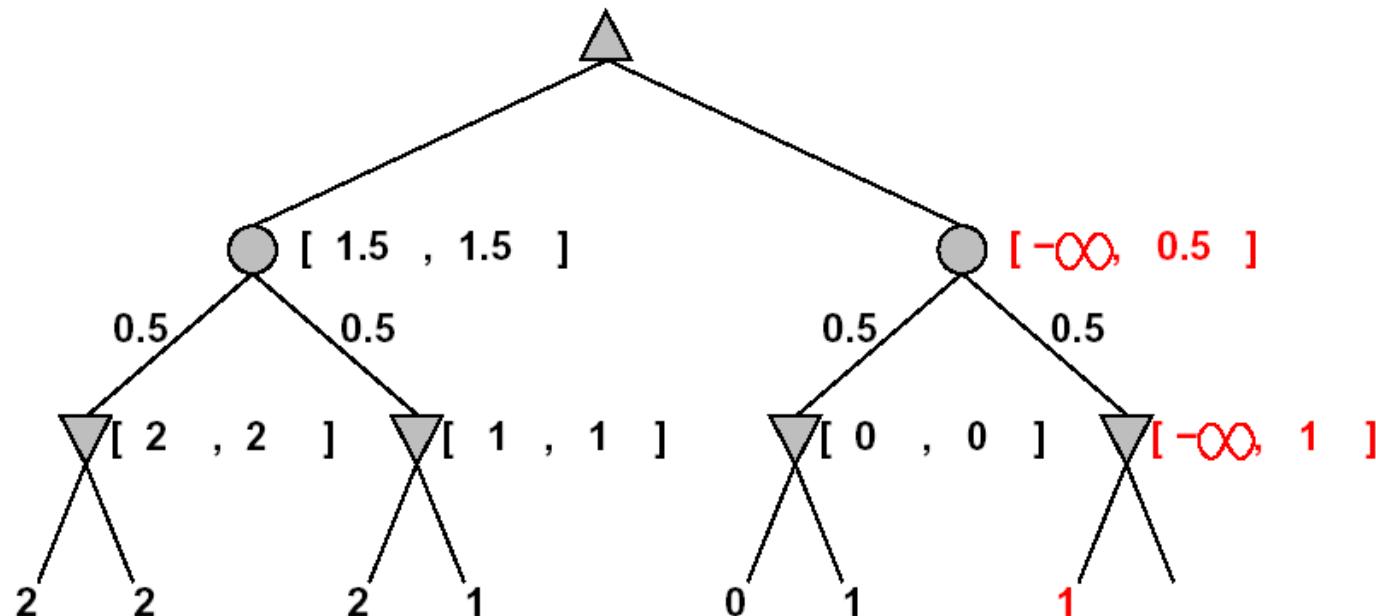
## Pruning in nondeterministic game trees

- A version of  $\alpha$ - $\beta$  pruning is possible:



## Pruning in nondeterministic game trees

- A version of  $\alpha$ - $\beta$  pruning is possible:

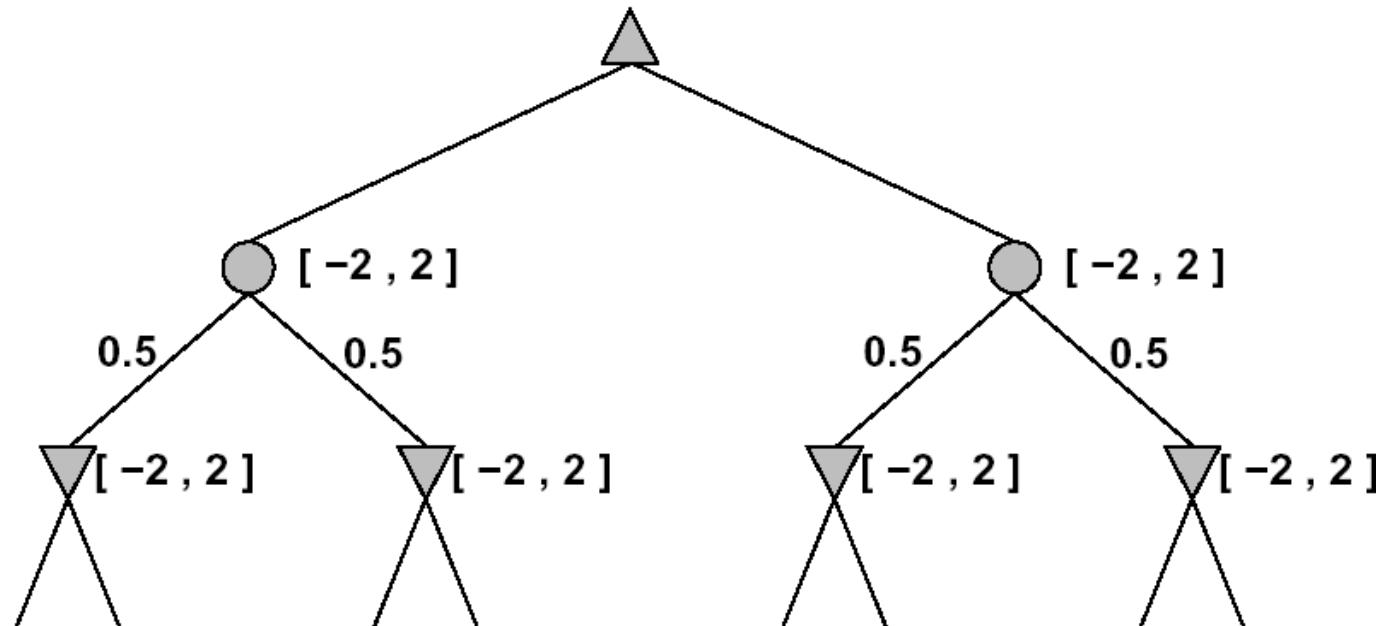




## Pruning (contd.)



- More pruning occurs if we can bound the leaf values

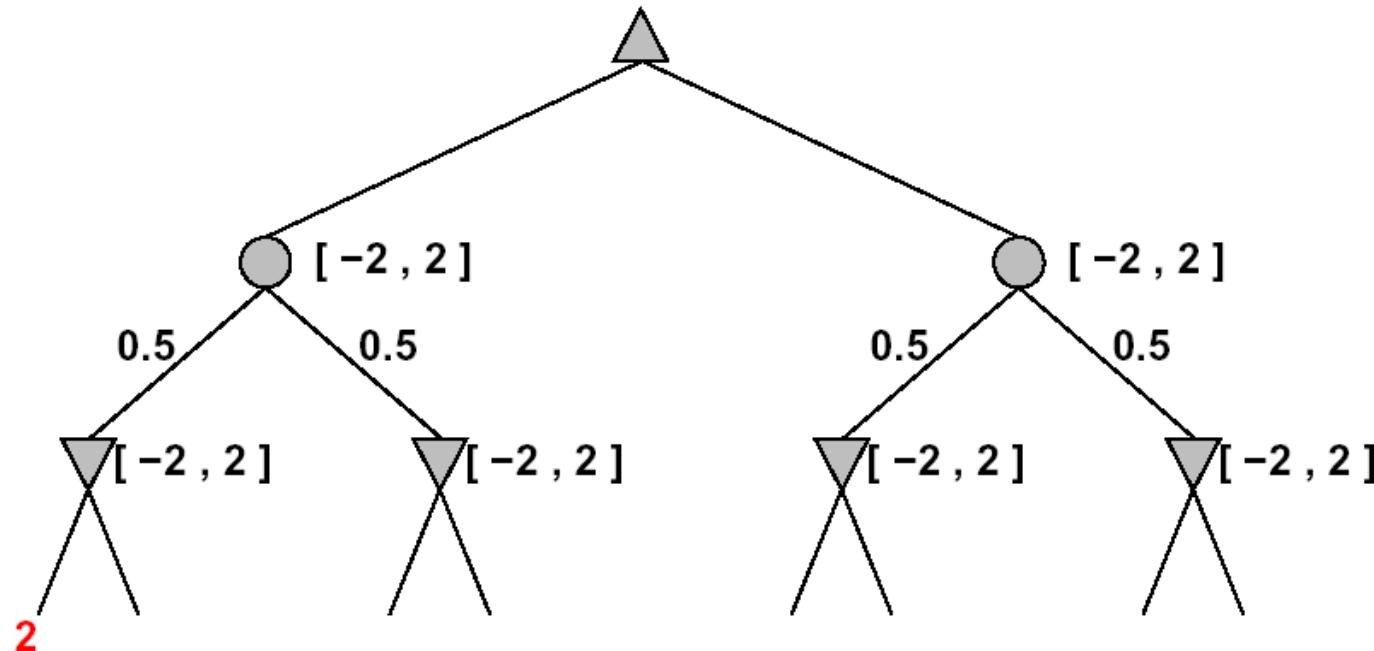




## Pruning (contd.)



- More pruning occurs if we can bound the leaf values

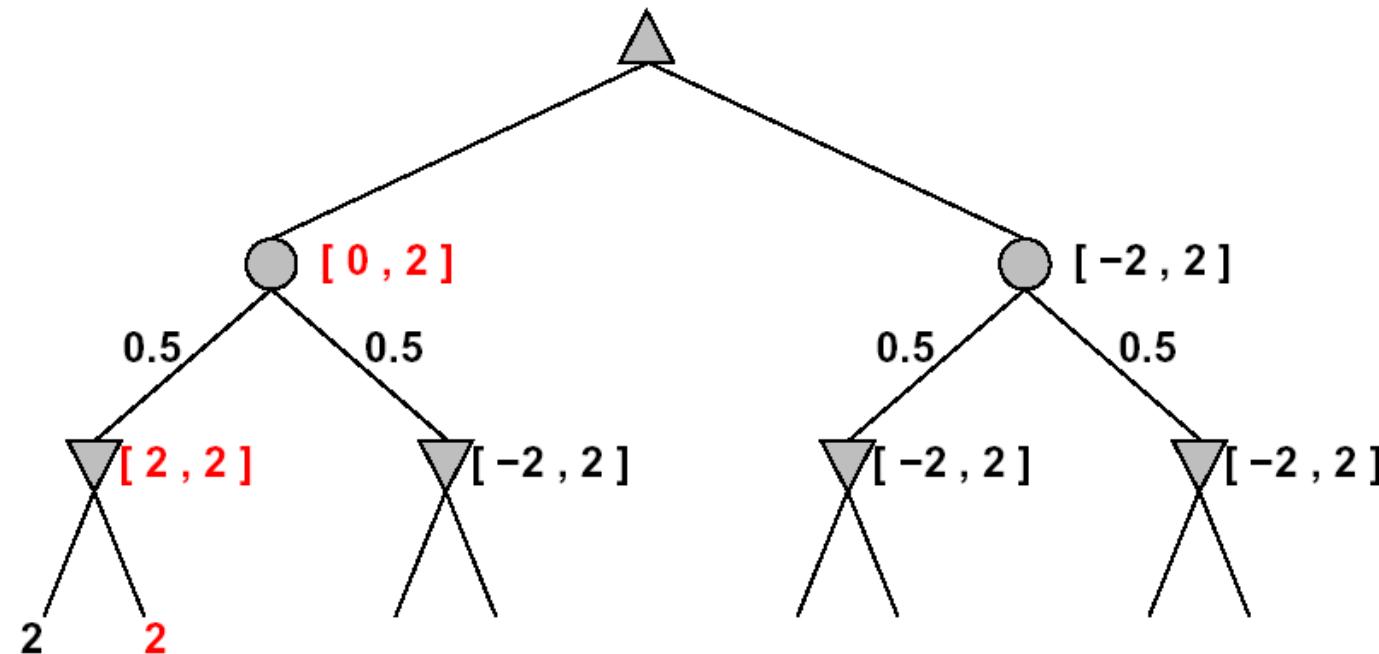




## Pruning (contd.)



- More pruning occurs if we can bound the leaf values

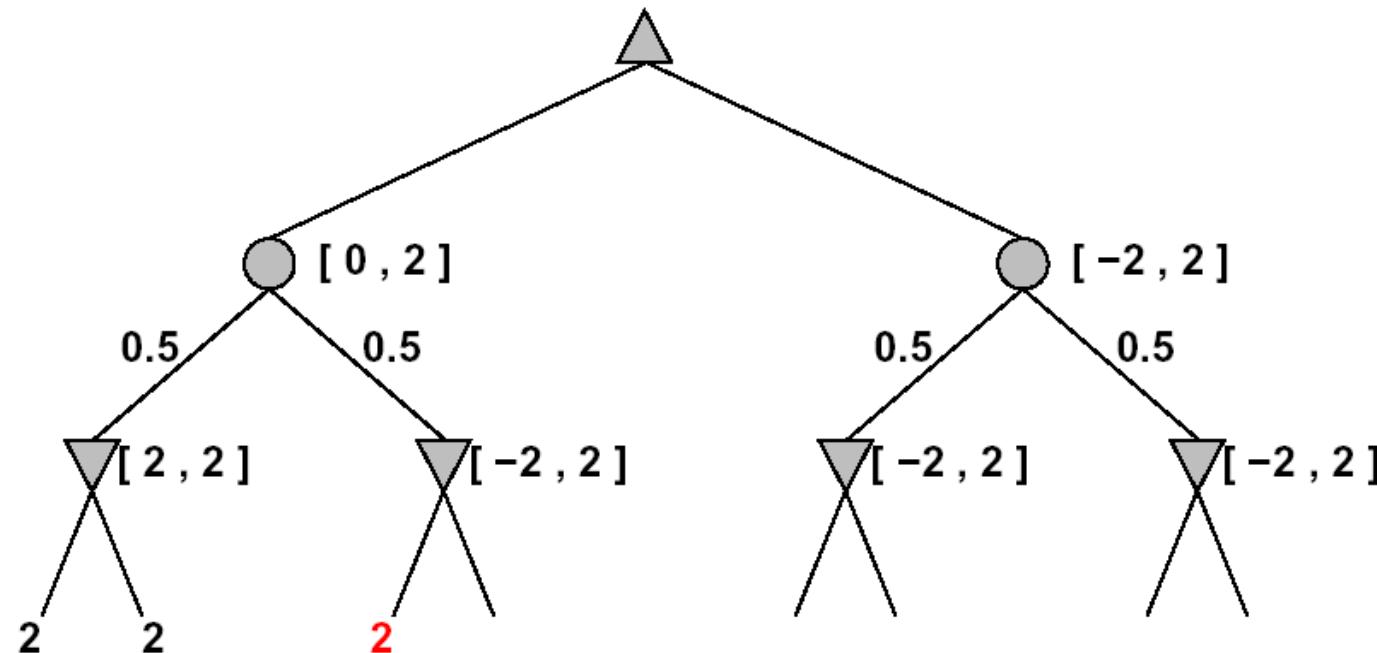




## Pruning (contd.)



- More pruning occurs if we can bound the leaf values

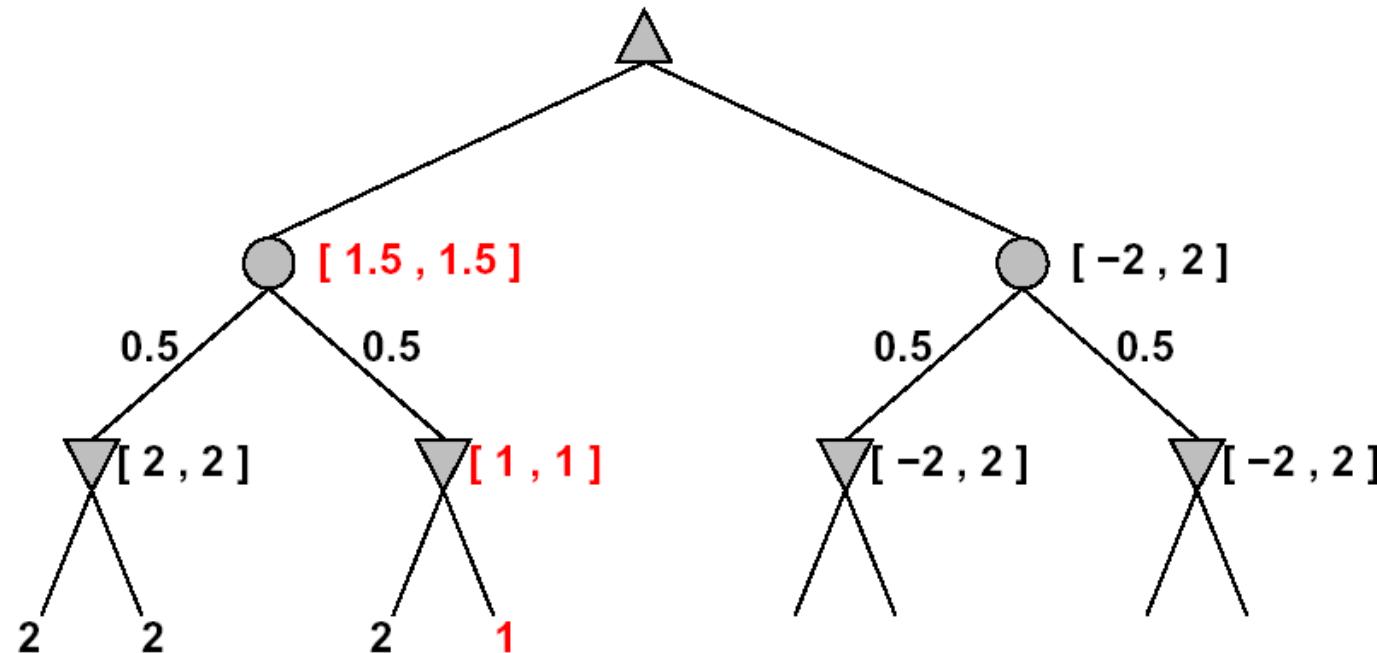




## Pruning (contd.)



- More pruning occurs if we can bound the leaf values

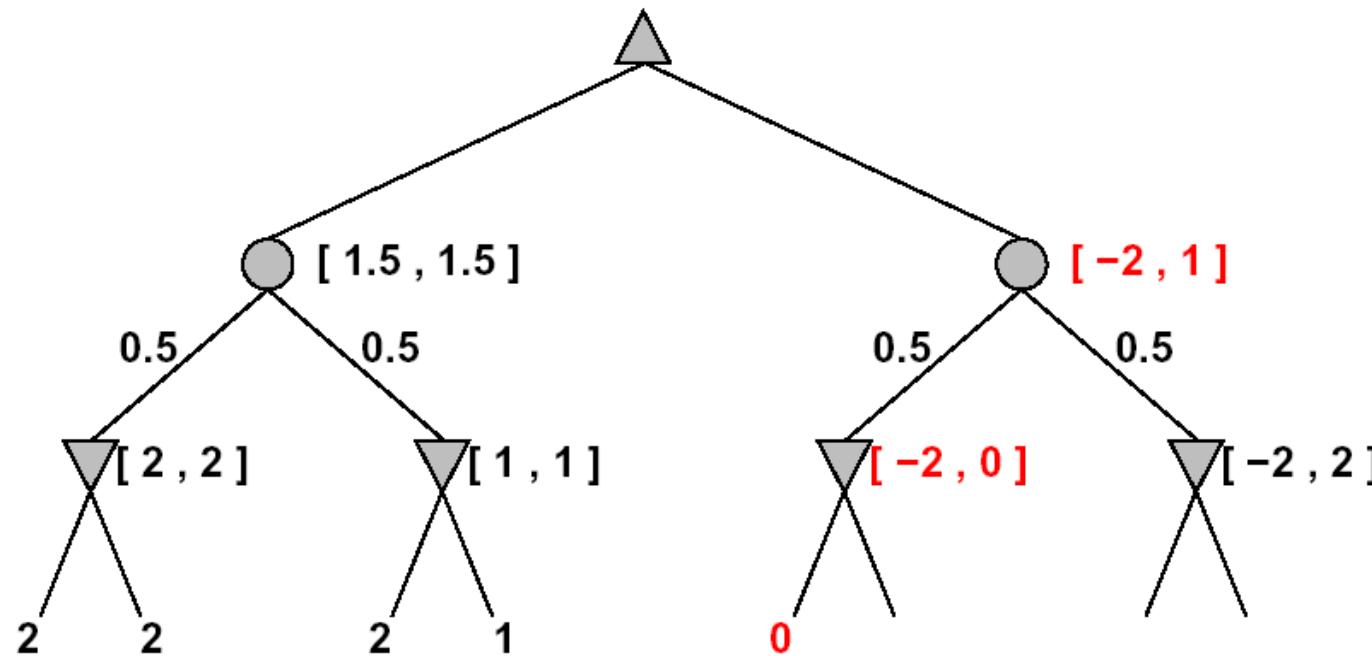




## Pruning (contd.)



- More pruning occurs if we can bound the leaf values

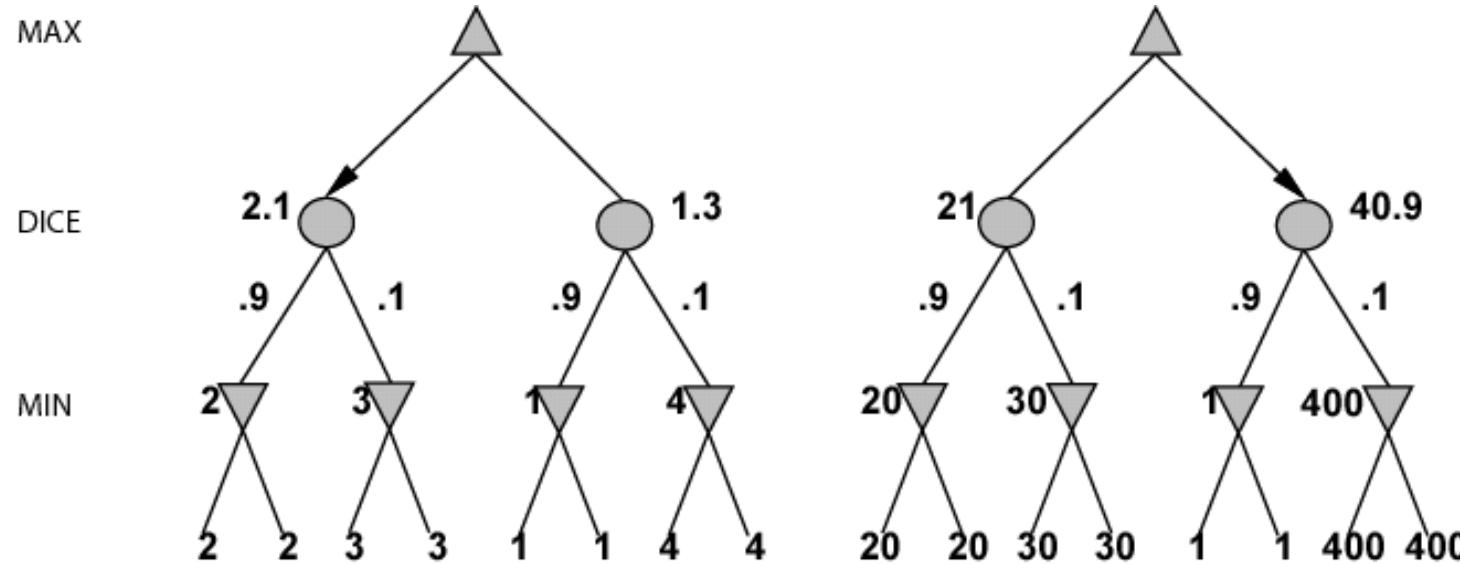




## Nondeterministic games in practice

- Dice rolls increase  $b$ : 21 possible rolls with 2 dice
- Backgammon  $\approx 20$  legal moves (can be 6,000 with 1-1 roll)  
 $\text{depth } 4 = 20 \ (21 \times 20)^3 \approx 1.2 \times 10^9$
- As depth increases, probability of reaching a given node shrinks  
 $\Rightarrow$  value of lookahead is diminished
- $\alpha-\beta$  pruning is much less effective
- **TDGammon** uses depth-2 search + very good **Eval**  
 $\approx$  world-champion level

## └ Digression:   └ Exact values DO matter



- Behaviour is preserved only by *positive linear* transformation ofc **Eval**
- Hence **Eval** should be proportional to the expected payoff



# Games of imperfect information

- E.g., card games, where opponent's initial cards are unknown
- Typically we can calculate a probability for each possible deal
- Seems just like having one big dice roll at the beginning of the game\*

**Idea:** compute the minimax value of each action in each deal,  
then choose the action with highest expected value over all deals\*

- Special case: if an action is optimal for all deals, it's optimal.\*
- GIB, current best bridge program, approximates this idea by
  - 1) generating 100 deals consistent with bidding information
  - 2) picking the action that wins most tricks on average



## Summary



- Games are fun to work on!  
(and dangerous)
- They illustrate several important points about AI
  - perfection is unattainable  $\Rightarrow$  must approximate
  - good idea to think about what to think about
  - uncertainty constrains the assignment of values to states
- Games are to AI as grand prix racing is to automobile design