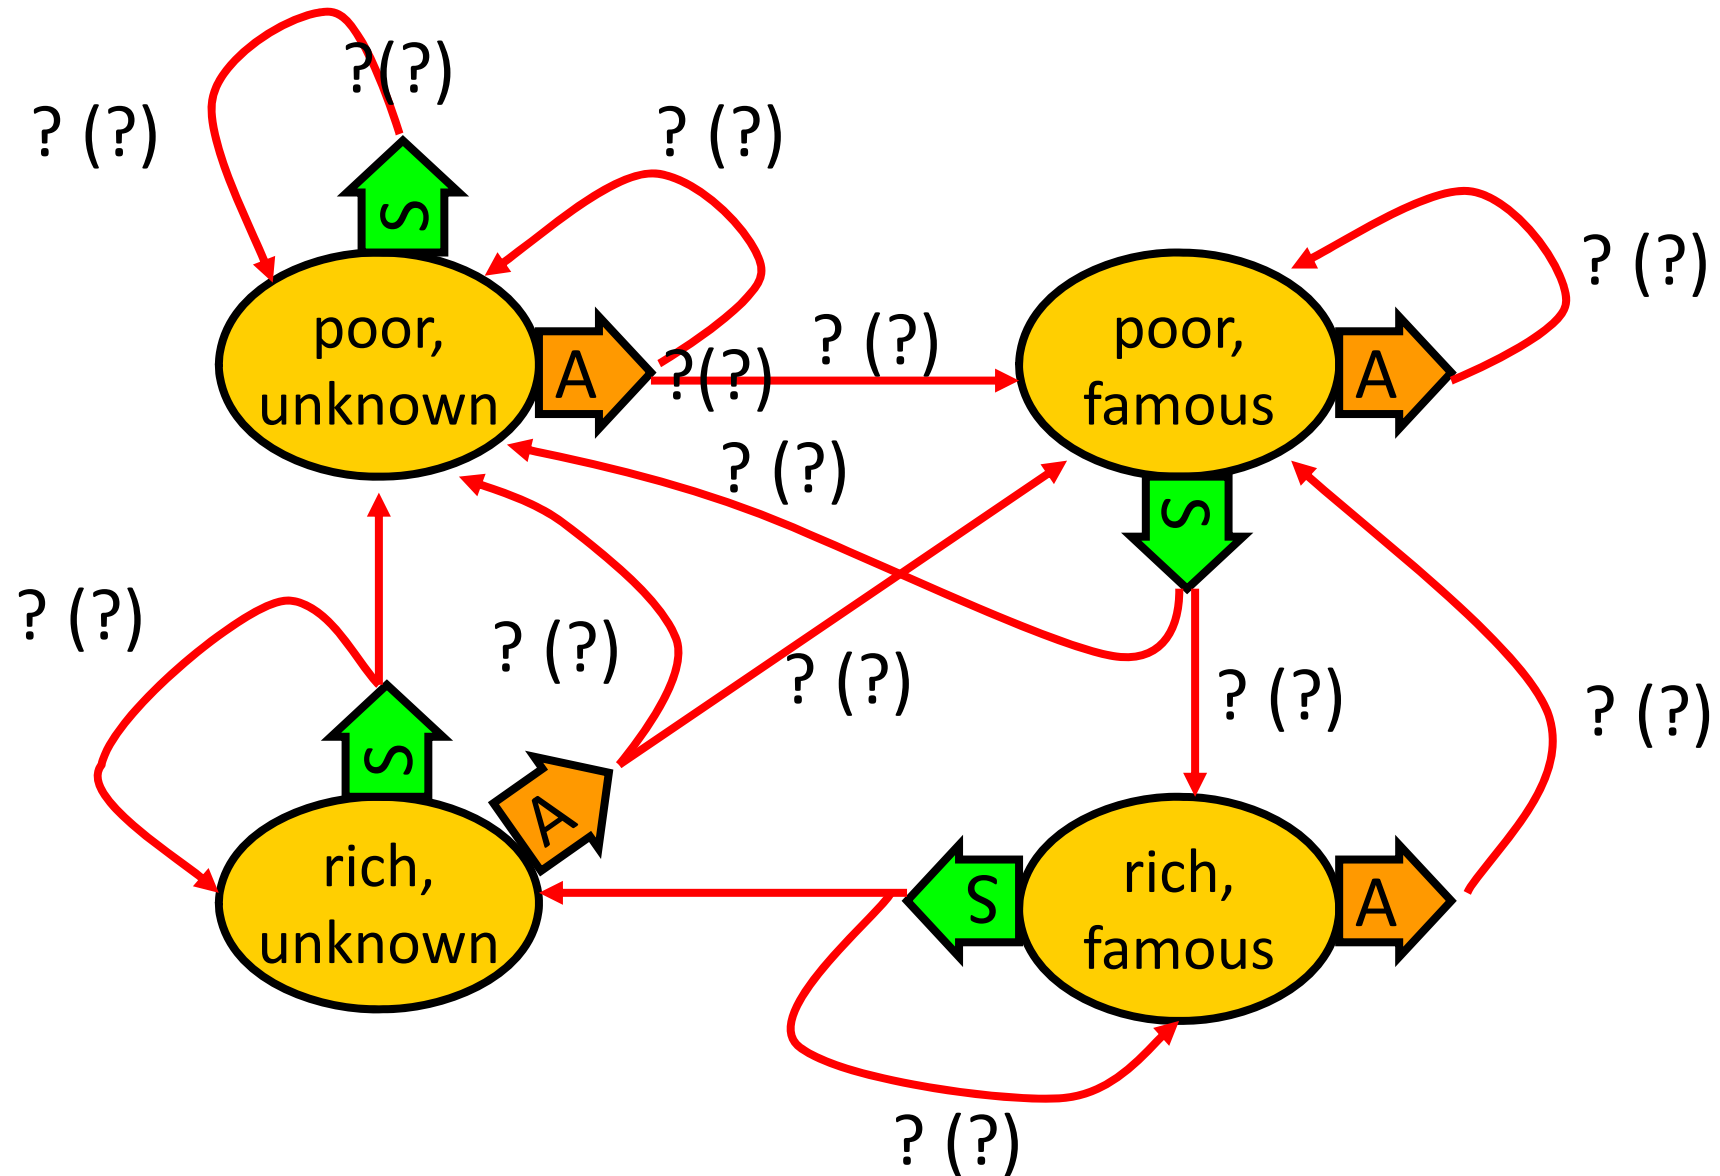


Probabilistic Foundations of Artificial Intelligence

Reinforcement Learning

Prof. Andreas Krause
Learning and Adaptive Systems (las.ethz.ch)

RL = Planning in unknown MDPs



Reinforcement Learning

- RL is different from supervised learning
 - The data we get is not i.i.d.
 - In reinforcement learning, the data we get *depends on our actions!*
 - Some actions have higher rewards than others!
- *Exploration—Exploitation Dilemma*: Should we
 - **Explore**: gather more data to avoid missing out on a potentially large reward?
 - **Exploit**: stick with our current knowledge and build an optimal policy for the data we've seen?

Possible approaches

- Always pick a random action?
 - Will eventually correctly estimate all probabilities and rewards 😊
 - May do extremely poorly! 😞
- Always pick the best action according to current knowledge?
 - Quickly get some reward
 - Can get stuck in suboptimal action! 😞
- Balance exploration and exploitation (more later)

Two basic approaches

1) Model-based RL (“Approximate dynamic programming”)

- Learn the MDP

Estimate transition probabilities $P(x' \mid x, a)$

Estimate reward function $r(x, a)$

- Optimize policy based on estimated MDP

2) Model-free RL

- Estimate the value function directly;
- Actor-critic methods;
- Policy gradient methods

Learning the MDP

- Need to estimate

- transition probabilities $P(X_{t+1} | X_t, A)$
- Reward function $r(X, A)$

- Can use techniques from learning Bayes Nets:
(regularized) maximum likelihood estimation

- Data set: "Experiences" $x_1, a_1, r_1, x_2, a_2, r_2, x_3, \dots$
 $D = \{ (x_1, a_1, r_1, x_2), (x_2, a_2, r_2, x_3), \dots, (x_{t-1}, a_{t-1}, r_{t-1}, x_t) \}$

- Estimate transitions:

$$P(X_{t+1} | X_t, A) \approx \frac{\text{Count}(X_{t+1}, X_t, A)}{\text{Count}(X_t, A)}$$

- Estimate rewards:

$$r(x, a) \approx \frac{1}{N_{x,a}} \sum_{t: X_t=x, A_t=a} R_t$$

Encouraging Exploration

- ε_t greedy
 - With probability ε_t : Pick random action
 - With probability $(1-\varepsilon_t)$: Pick best action
- If sequence ε_t satisfies Robbins Monro (RM) conditions then will converge to optimal policy with probability 1

$\varepsilon_t \rightarrow 0$, in particular: $\sum_{t=1}^{\infty} \varepsilon_t \rightarrow \infty$

F.g. $\varepsilon_t = \frac{1}{t}$
 $\varepsilon_t = \min\left(1, \frac{C}{t}\right)$

$\sum_{t=1}^{\infty} \varepsilon_t^2$ finite


- Often performs quite well
- Doesn't quickly eliminate clearly suboptimal actions

The R_{\max} Algorithm [Brafman & Tenenholz '02]


Optimism in the face of uncertainty!

- If you don't know $r(x, a)$:
 - Set it to R_{\max} !
- If you don't know $P(x' | x, a)$:
 - Set $P(x^* | x, a) = 1$ where x^* is a “**fairy tale**” state:

$$P(x^* | x^*, a) = 1 \quad \forall a$$
$$r(x^*, a) = R_{\max}$$



Implicit Exploration Exploitation in R_{\max}



$r(1, \text{Dig})=0$ $r(2, \text{Dig})=0$

x

Three actions:

- Left
- Right
- Dig

$$r(i, \text{Left}) = 0$$

$$r(i, \text{Right}) = 0$$

Never need to explicitly choose whether we're exploring or exploiting!

Can rule out clearly suboptimal actions very quickly

Exploration—Exploitation Lemma

Theorem: Every T timesteps, w.h.p., R_{\max} either

- Obtains near-optimal reward, or
 - Visits at least one unknown state-action pair
-
- T is related to the mixing time of the Markov chain of the MDP induced by the optimal policy

The R_{\max} algorithm

Input: Starting state x_0 , discount factor γ

Initially:

- Add fairy tale state x^* to MDP
- Set $r(x, a) = R_{\max}$ for all states x and actions a
- Set $P(x^* | x, a) = 1$ for all states x and actions a
- Choose optimal policy for r and P

Repeat:

- Execute policy π
- For each visited state action pair x, a , update $r(x, a)$
- Estimate transition probabilities $P(x' | x, a)$
- If observed “enough” transitions / rewards, recompute policy π according to current model P and r

How much is “enough”?

How many samples do we need to accurately estimate $P(x' | x, a)$ or $r(x, a)$?

Hoeffding-Chernoff bound:

- X_1, \dots, X_n i.i.d. samples from Bernoulli distribution w. mean μ

$$P\left(\left|\mu - \frac{1}{n} \sum_i X_i\right| \geq \varepsilon\right) \leq 2 \exp(-2n\varepsilon^2)$$

Performance of R_{\max} [Brafman & Tenenholz]

Theorem:

With probability $1-\delta$, R_{\max} will reach an ε -optimal policy in a number of steps that is polynomial in $|X|$, $|A|$, T , $1/\varepsilon$ and $\log(1/\delta)$

Problems of model-based RL?

- Memory required: "Dense" MDPs, $O(|X|^2 \cdot |A|)$ to store frequencies / rewards
(in general $O(\# \text{ edges in MDP})$)
- Computation time:
need to replan "often"
each replanning step is poly $(|X|, |A|)$

Two basic approaches

1) Model-based RL (“Approximate dynamic programming”)

- Learn the MDP

Estimate transition probabilities $P(x' \mid x, a)$

Estimate reward function $r(x, a)$

- Optimize policy based on estimated MDP

2) Model-free RL

- Estimate the value function directly;
- Actor-critic methods;
- Policy gradient methods

Model free RL

- Recall:

1. Optimal value function $V^*(x) \rightarrow$ opt. policy π^*
2. For optimal value function it holds:

$$V^*(x) = \max_a Q^*(x, a)$$

$$\text{where } Q^*(x, a) = r(x, a) + \gamma \sum_{x'} P(x' \mid x, a) V^*(x')$$

Key idea: Estimate $Q^*(x, a)$ directly from samples!

Q-learning

- Estimate $Q^*(x, a) = r(x, a) + \gamma \sum_{x'} P(x' | x, a) V^*(x')$

$$V^*(x) = \max_a Q^*(x, a)$$

Expected value of next state

- Suppose we

- Have initial estimate of $Q(x, a)$
- observe transition x, a, x' with reward r

↳ approx. by
 - old est. of V^*
 - single sample
 $x' \sim P(x' | x, a)$

$$Q(x, a) \leftarrow (1 - \alpha_t) \underbrace{Q^{\text{old}}(x, a)}_{\text{old estimate for } Q} + \alpha_t \left(r + \gamma \underbrace{\max_{a'} Q^{\text{old}}(x', a')}_{\text{old estimate for } V(x')} \right)$$

"new estimate for Q
 given single exp. only"

Q-learning

$$Q(x, a) \leftarrow (1 - \alpha_t)Q(x, a) + \alpha_t \left(r + \gamma \max_{a'} Q(x', a') \right)$$

Theorem: If learning rate α_t satisfies

$$\sum_t \alpha_t = \infty$$

$$\sum_t \alpha_t^2 < \infty$$

and actions are chosen at random*, then Q learning converges to optimal Q^* with probability 1

How can we trade off exploration and exploitation?

Convergence of Optimistic Q-learning

[Even-dar & Mansour '02]

Similar to R_{\max} :

Initialize $Q(x, a) = \frac{R_{\max}}{1 - \gamma} \prod_{t=1}^{T_{\text{init}}} (1 - \alpha_t)^{-1}$

Theorem: With prob. $1 - \delta$, optimistic Q -learning obtains an ε -optimal policy after a number of time steps that is polynomial in $|X|$, $|A|$, $1/\varepsilon$ and $\log(1/\delta)$

Properties of Q-learning

- Memory required: $O(|X| \cdot |A|)$ to store Q fn
- Computation time: $O(|A|)$

Acknowledgments

- Slides based on material accompanying the textbook “AI: A Modern Approach” (3rd edition) by S. Russell and P. Norvig, the textbook “Reinforcement Learning: An Introduction” by R. S. Sutton and A. G. Barto